Recurrent Neural Networks

Antoine Bosselut

Announcements

- **Assignment 1 Released:** Due **March 17, 2024**

  - **Q&A Sessions:**

    ‣ Wednesday, March 6th, 2024 - 1 PM (STCC)

    ‣ Thursday, March 14th, 2024 - 1 PM (CE 1 6)

- **Project Update:**

  - **Groups of 4 are allowed.**

  - There will however be an added component for Groups of 4,

  - Your project output will be evaluated on multiple languages rather than only English.

Some slides adapted from Danqi Chen, Mohit Iyyer, Dan Jurafsky

Section Outline

- **Fixing the context bottleneck:** recurrent neural networks

- **Training recurrent neural networks:** backpropagation through time

- **Training Challenges:** Vanishing Gradients

- **Mitigations:** LSTMs, GRUs

Sequences are not of consistent length

The **cat** chased the **mouse**

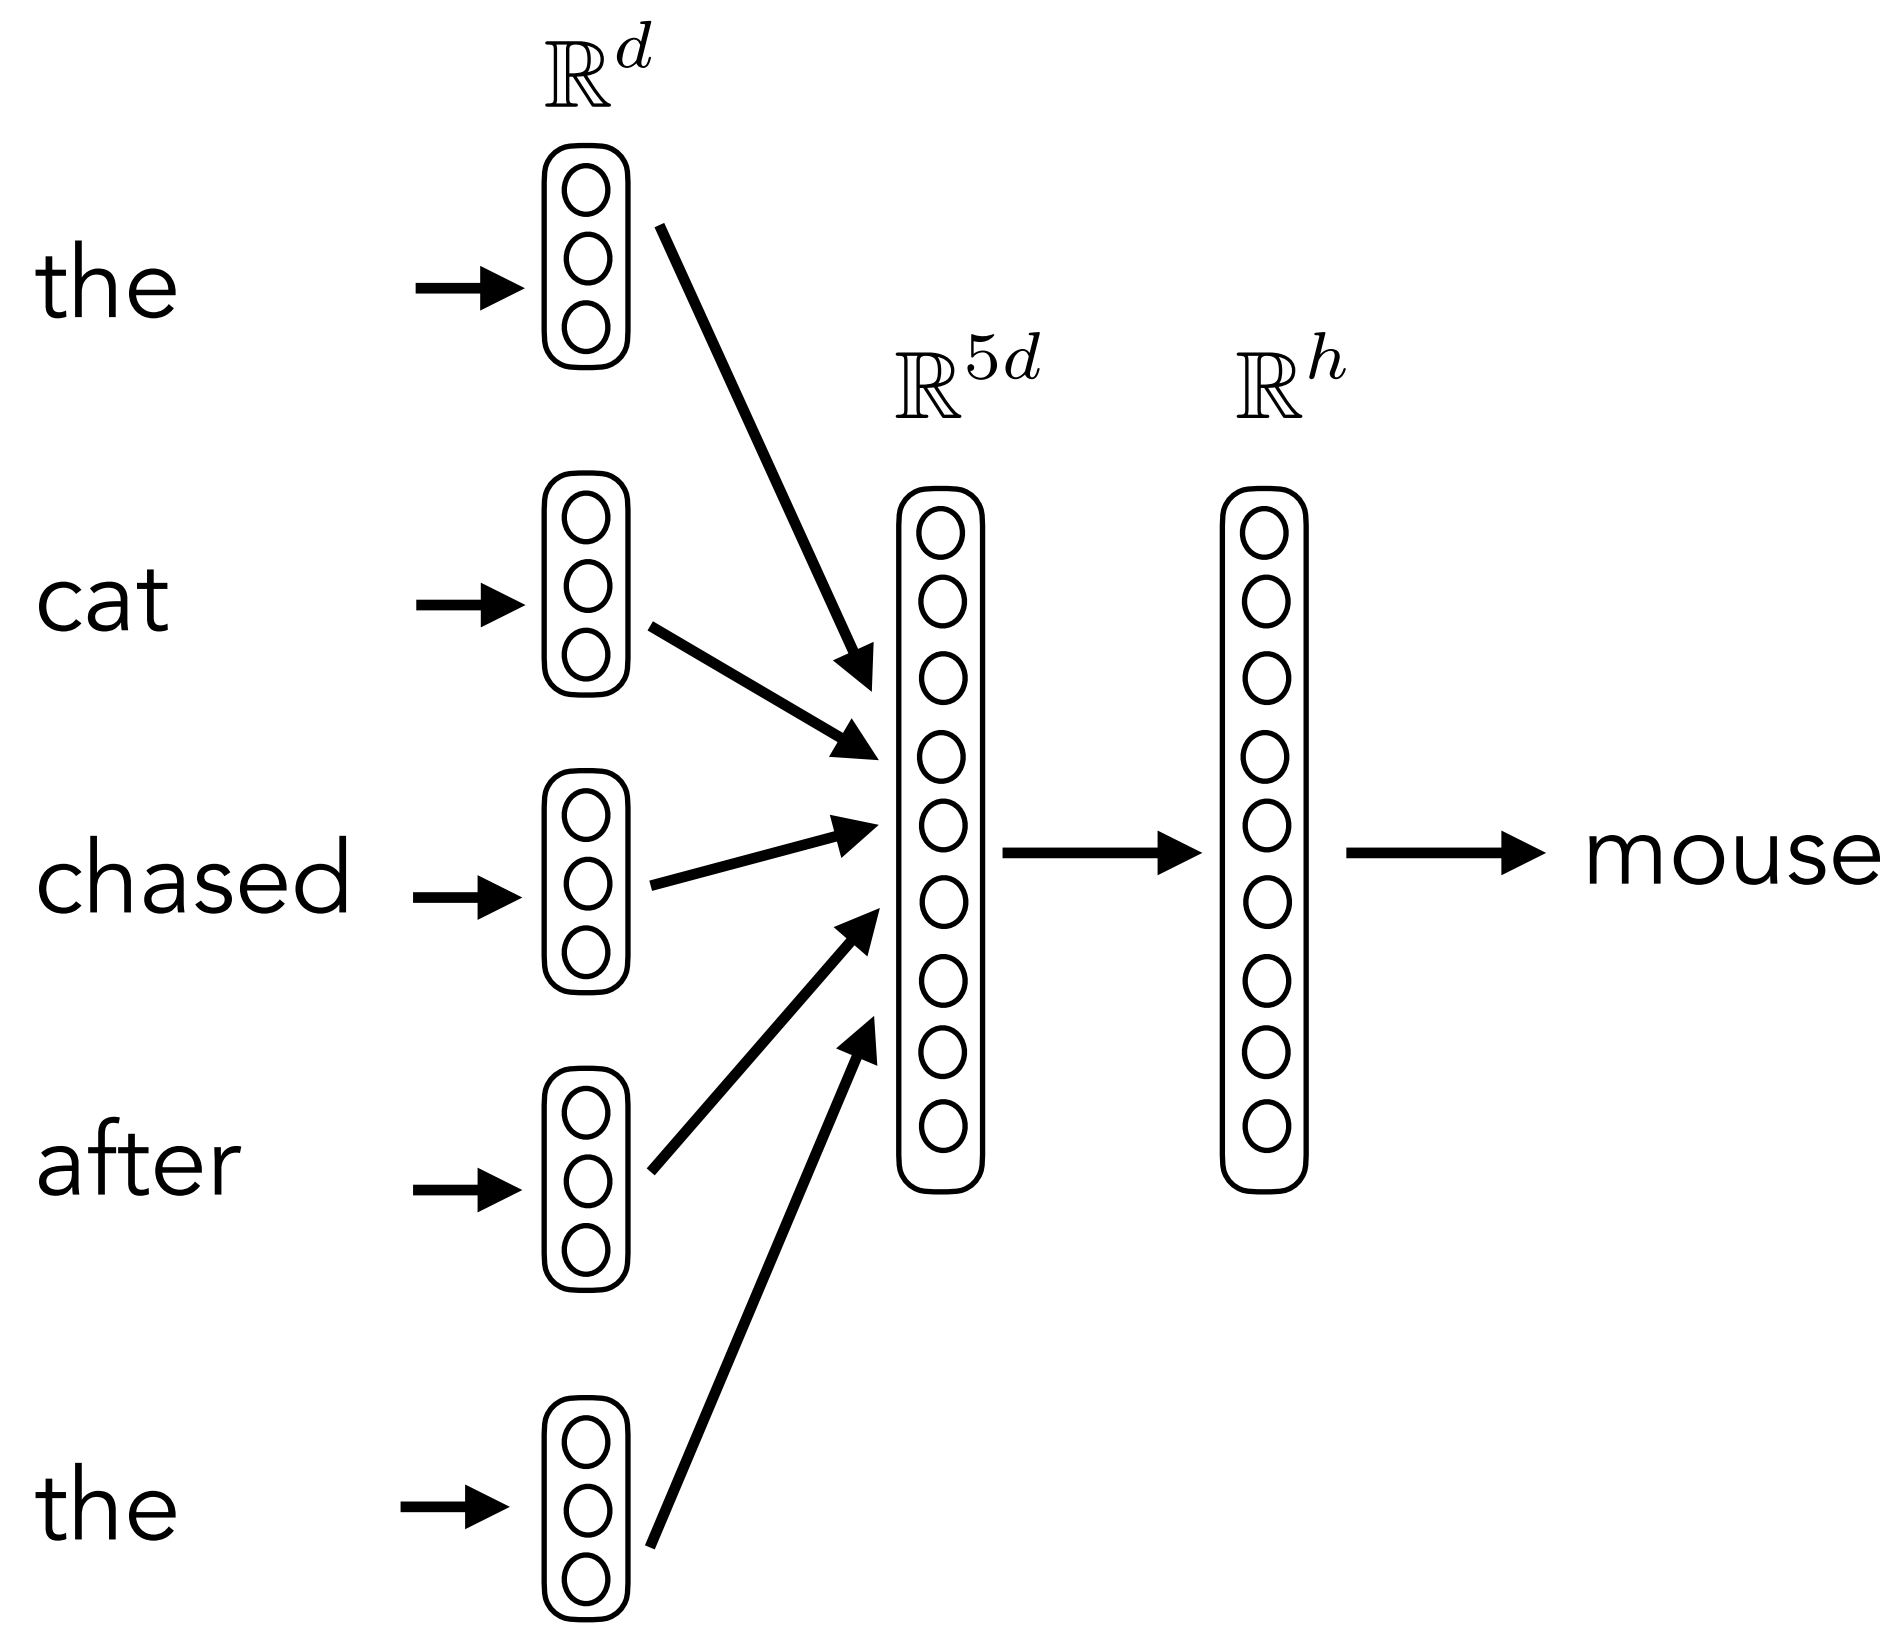The starving **cat** chased the **mouse**

The starving **cat** fanatically chased the **mouse**

The starving **cat** fanatically chased the elusive **mouse**

The starving **cat**, who had not eaten in six days, fanatically chased the elusive **mouse**

Fixed-context Neural Language Models

the $\rightarrow$ $\mathbb{R}^d$

cat $\rightarrow$

chased $\rightarrow$ $\mathbb{R}^{5d}$ $\mathbb{R}^h$ $\rightarrow$ mouse

after $\rightarrow$

the $\rightarrow$

P(mouse | the cat chased the) ✔

P(mouse | the starving cat chased the) ✔

P(mouse | starving cat chased after the) ✔

P(mouse | cat fanatically chased after the) ✔

P(mouse | fanatically chased after the elusive) ✘

5   (Bengio et al., 2003): A Neural Probabilistic Language Model

Recurrent Neural Networks

- **Solution:** Recurrent neural networks — NNs with feedback loops

Output
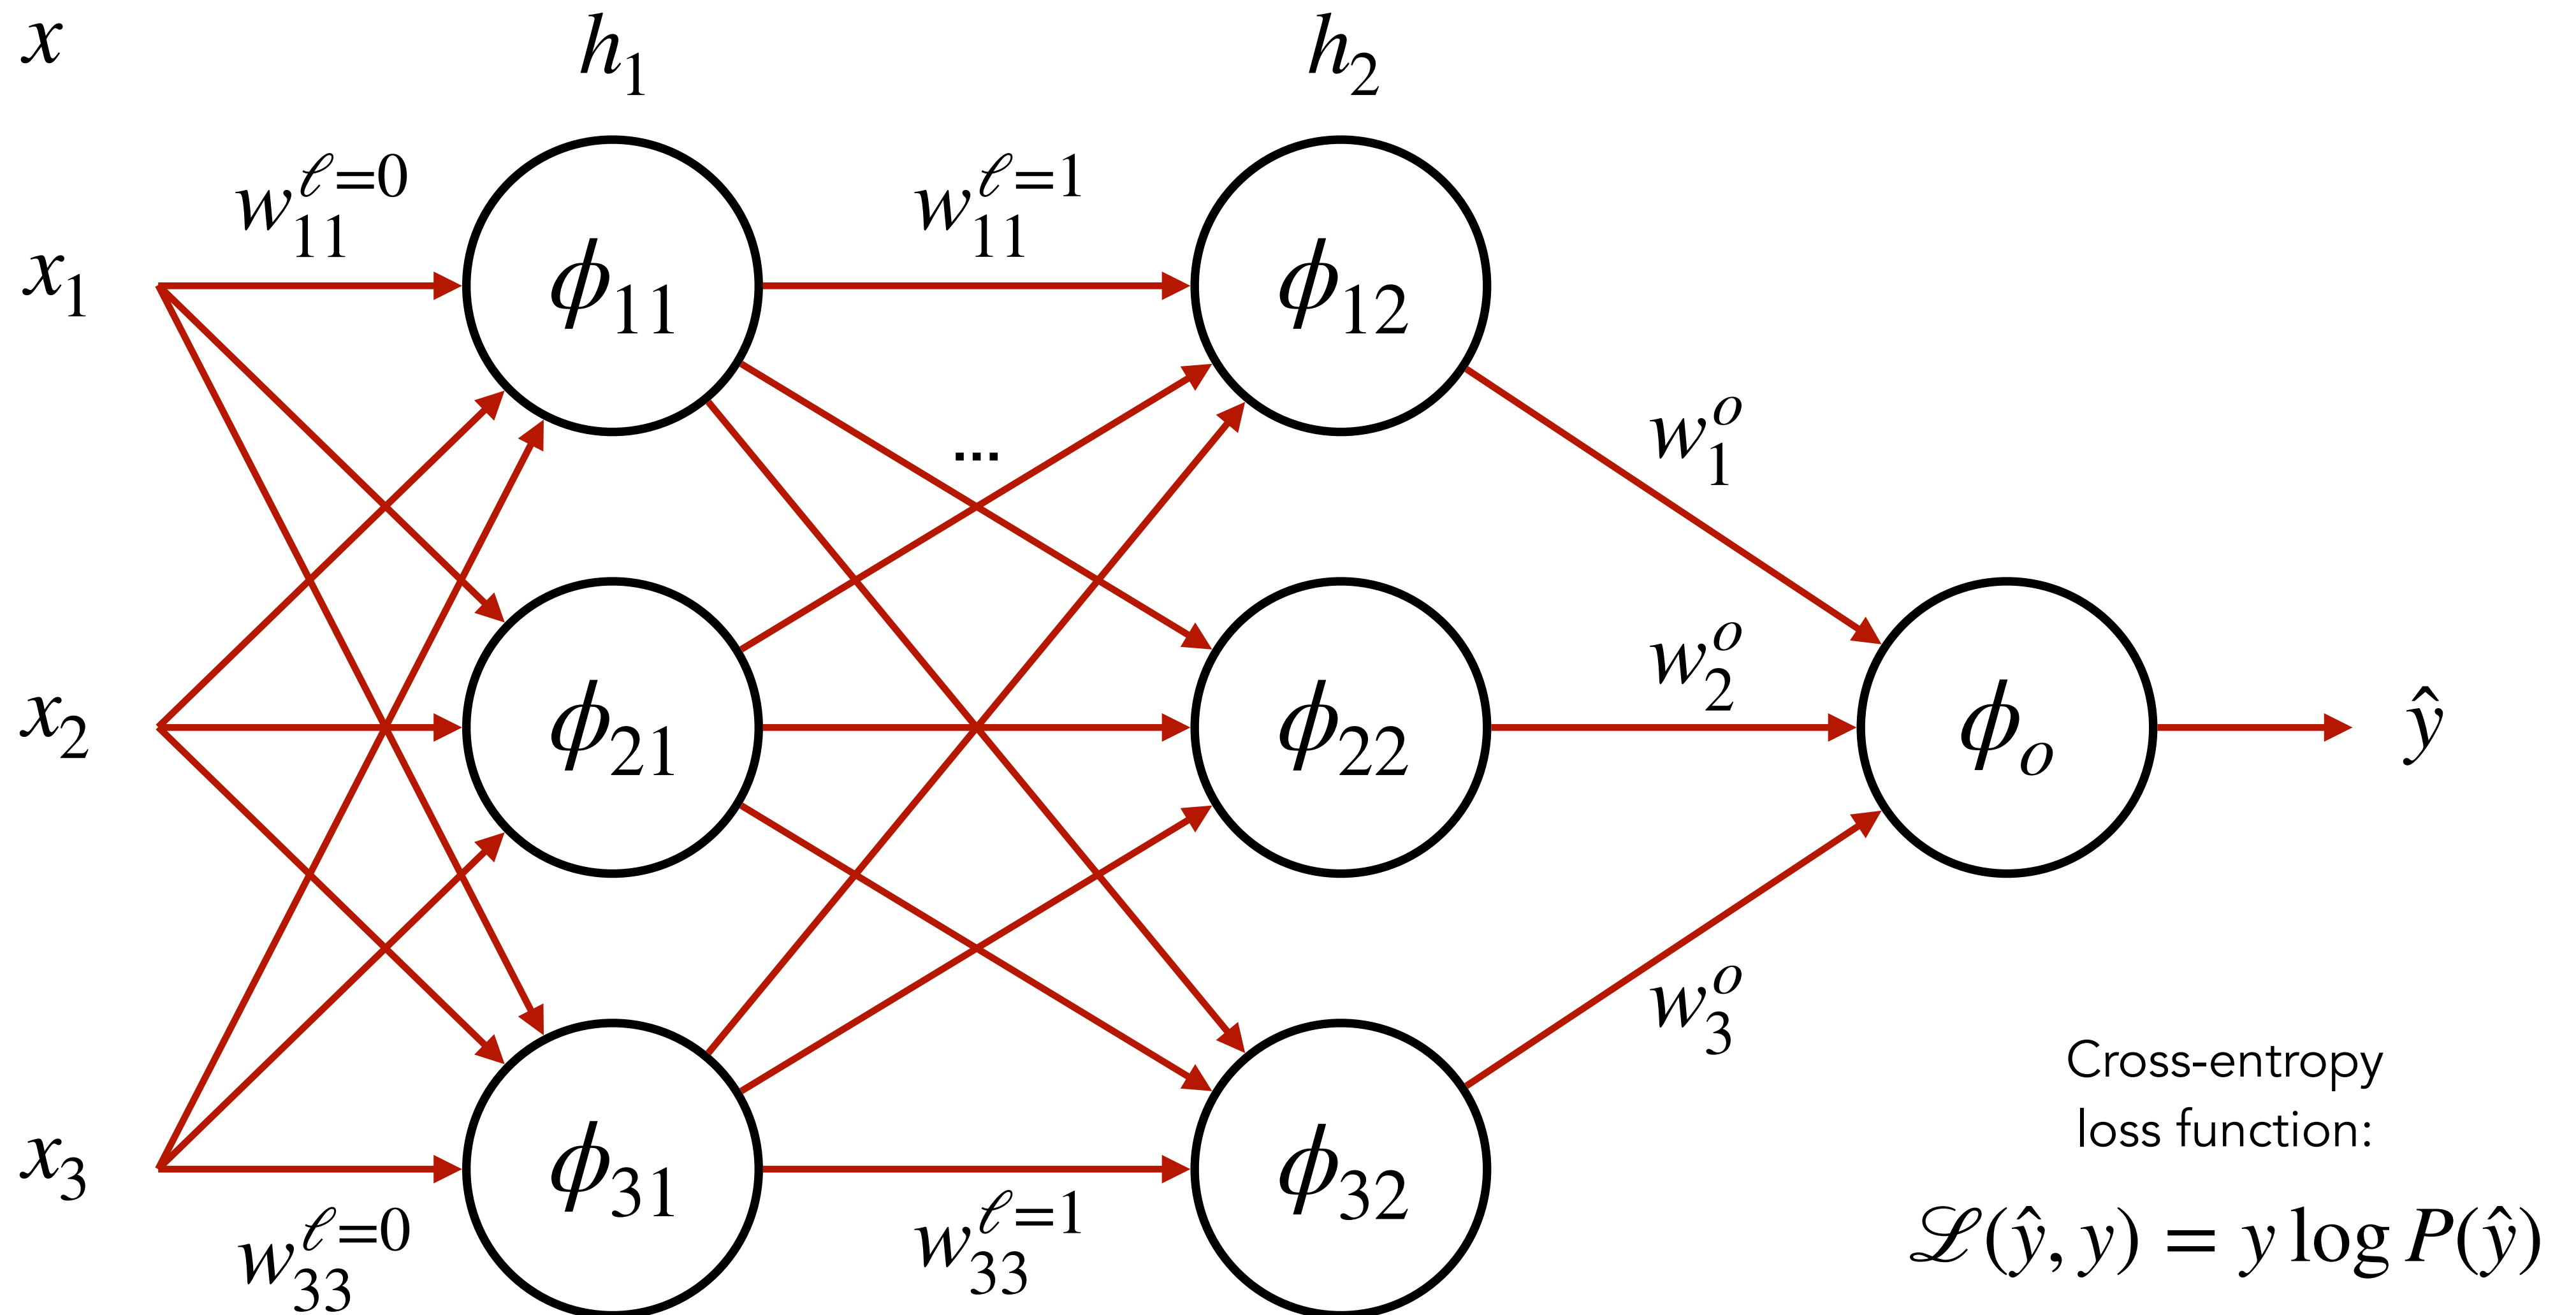
State

Input

h_t

z_t

x_t

Recurrent Neural Networks

Unrolling the RNN across all time steps gives full computation graph

Allows for learning from entire sequence history, regardless of length

Backpropagation Review: FFNs

x h_1 h_2

$w_{11}^{\ell=0}$ $w_{11}^{\ell=1}$

x_1 ϕ_{11} ϕ_{12} w_1^o

...

x_2 ϕ_{21} ϕ_{22} w_2^o ϕ_o $\hat{y}$

w_3^o

x_3 ϕ_{31} ϕ_{32}

$w_{33}^{\ell=0}$ $w_{33}^{\ell=1}$

Cross-entropy loss function:

$$\mathcal{L}(\hat{y}, y) = y \log P(\hat{y})$$

Backpropagation Review: FFNs

h_2

ϕ_{12}

ϕ_{22}

ϕ_{32}

w_1^o

w_2^o

w_3^o

ϕ_o

$\hat{y}$

$$\mathcal{L}(\hat{y}, y) = y \log P(\hat{y})$$

$$\hat{y} = \phi_o(u)$$

$$u = w_1^o \times \phi_{12}(\,.\,) + w_2^o \times \phi_{22}(\,.\,) + w_3^o \times \phi_{32}(\,.\,)$$

$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{12}(\,.\,)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(\,.\,)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^0$$

Backpropagation Review: FFNs

h_2

$\mathscr{L}(\hat{y}, y) = y \log P(\hat{y})$

$\hat{y} = \phi_o(u)$

$u = w_1^o \times \phi_{12}(\,.\,) + w_2^o \times \phi_{22}(\,.\,) + w_3^o \times \phi_{32}(\,.\,)$

$\dfrac{\partial \mathscr{L}(\hat{y}, y)}{\partial \phi_{12}(\,.\,)} = \dfrac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \dfrac{\partial \hat{y}}{\partial u} \dfrac{\partial u}{\partial \phi_{12}(\,.\,)}$

$= \dfrac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \dfrac{\partial \phi_o(u)}{\partial u} w_1^0$

Backpropagation Review: FFNs

h_2

ϕ_{12}

ϕ_{22}

ϕ_{32}

w_1^o

w_2^o

w_3^o

ϕ_o

$\hat{y}$

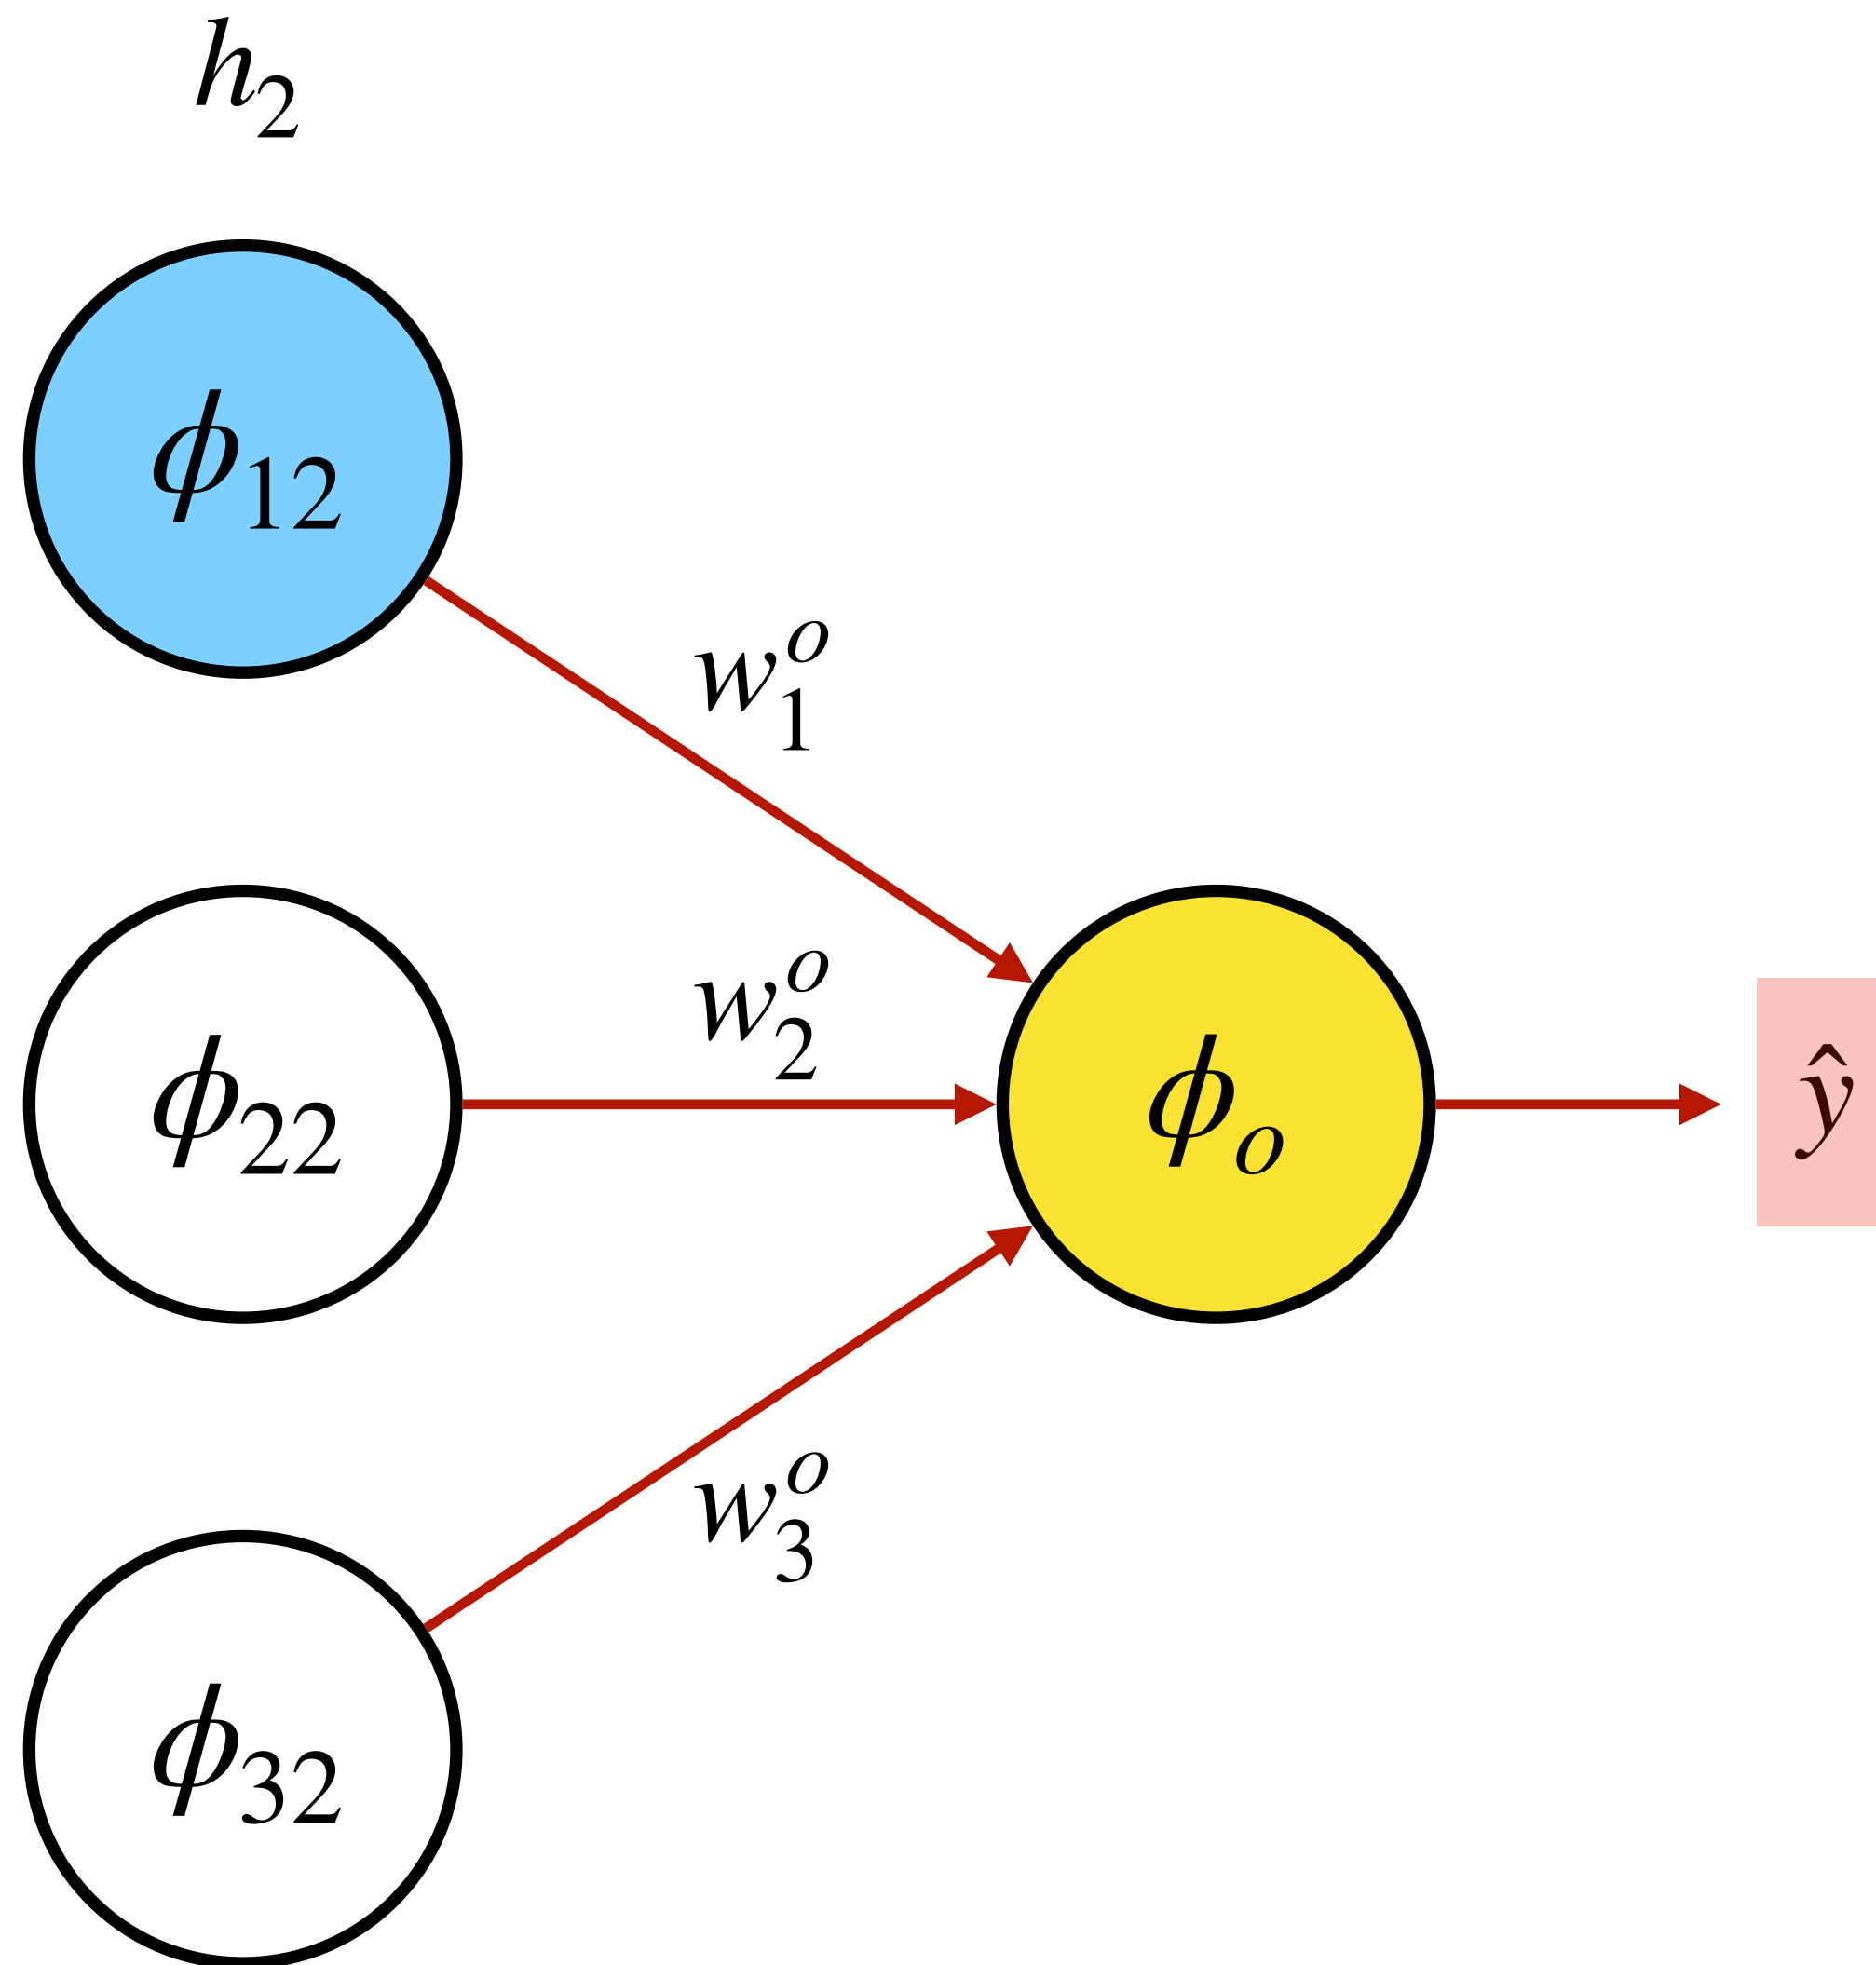$\mathscr{L}(\hat{y}, y) = y \log P(\hat{y})$

$\hat{y} = \phi_o(u)$

$u = w_1^o \times \phi_{12}(\,.\,) + w_2^o \times \phi_{22}(\,.\,) + w_3^o \times \phi_{32}(\,.\,)$

$$\frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \phi_{12}(\,.\,)} = \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(\,.\,)}$$

$$= \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^0$$

Depends on label y

Backpropagation Review: FFNs

h_2

$\mathscr{L}(\hat{y}, y) = \boxed{y \log P(\hat{y})} + \boxed{(1-y)\log P(1-\hat{y})}$

$\hat{y} = \phi_o(u)$

$u = \boxed{w_1^o \times \phi_{12}(\,.\,)} + w_2^o \times \phi_{22}(\,.\,) + w_3^o \times \phi_{32}(\,.\,)$

$\dfrac{\partial \mathscr{L}(\hat{y}, y)}{\partial \phi_{12}(\,.\,)} = \dfrac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \dfrac{\partial \hat{y}}{\partial u} \boxed{\dfrac{\partial u}{\partial \phi_{12}(\,.\,)}}$

$= \boxed{\dfrac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}}} \boxed{\dfrac{\partial \phi_o(u)}{\partial u}} \boxed{w_1^0}$

Depends on label y

Depends on ϕ_o

Backpropagation Review: FFNs

h_2

$$\mathscr{L}(\hat{y}, y) = \boxed{y \log P(\hat{y})} + \boxed{(1-y)\log P(1-\hat{y})}$$

$$\hat{y} = \phi_o(u)$$

$$u = w_1^o \times \phi_{12}(.) + \boxed{w_2^o \times \phi_{22}(.)} + w_3^o \times \phi_{32}(.)$$

$$\frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \phi_{22}(.)} = \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \boxed{\frac{\partial u}{\partial \phi_{22}(.)}}$$

$$= \boxed{\frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}}} \boxed{\frac{\partial \phi_o(u)}{\partial u}} \boxed{w_2^0}$$

Depends on label y

Depends on ϕ_o

Backpropagation Review: FFNs

h_1

h_2

ϕ_{11}

ϕ_{12}

$w_{11}^{\ell=1}$

ϕ_{21}

ϕ_{22}

w_1^o

ϕ_{31}

ϕ_{32}

$w_{33}^{\ell=1}$

w_2^o

w_3^o

ϕ_o

$\hat{y}$

...

$$\frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \phi_{12}(.)} = \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(.)}$$

$$= \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^0$$

$$v = w_{11}^{\ell=1} \times \phi_{11}(.) + w_{21}^{\ell=1} \times \phi_{21}(.) + w_{31}^{\ell=1} \times \phi_{31}(.)$$

$$\frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \phi_{11}(.)} = \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(v)} \frac{\partial \phi_{12}(v)}{\partial v} \frac{\partial v}{\partial \phi_{11}(.)}$$

$$= \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^0 \frac{\partial \phi_{12}(v)}{\partial v} w_{11}^{\ell=1}$$

Backpropagation Review: FFNs

h_1

h_2

ϕ_{11}

$w_{11}^{\ell=1}$

ϕ_{12}

w_1^o

...

ϕ_{21}

ϕ_{22}

w_2^o

ϕ_o

$\hat{y}$

ϕ_{31}

$w_{33}^{\ell=1}$

ϕ_{32}

w_3^o

$$\frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \phi_{12}(.)} = \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(.)}$$

$$= \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^0$$

$$v = w_{11}^{\ell=1} \times \phi_{11}(.) + w_{21}^{\ell=1} \times \phi_{21}(.) + w_{31}^{\ell=1} \times \phi_{31}(.)$$

$$\frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \phi_{11}(.)} = \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(v)} \frac{\partial \phi_{12}(v)}{\partial v} \frac{\partial v}{\partial \phi_{11}(.)}$$

$$= \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^0 \frac{\partial \phi_{12}(v)}{\partial v} w_{11}^{\ell=1}$$

Backpropagation Review: FFNs

h_1 h_2

ϕ_{11} $w_{11}^{\ell=1}$ ϕ_{12}

ϕ_{21} ϕ_{22} w_2^o ϕ_o $\hat{y}$

w_1^o

w_3^o

ϕ_{31} $w_{33}^{\ell=1}$ ϕ_{32}

$$\frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \phi_{12}(.)} = \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(.)}$$

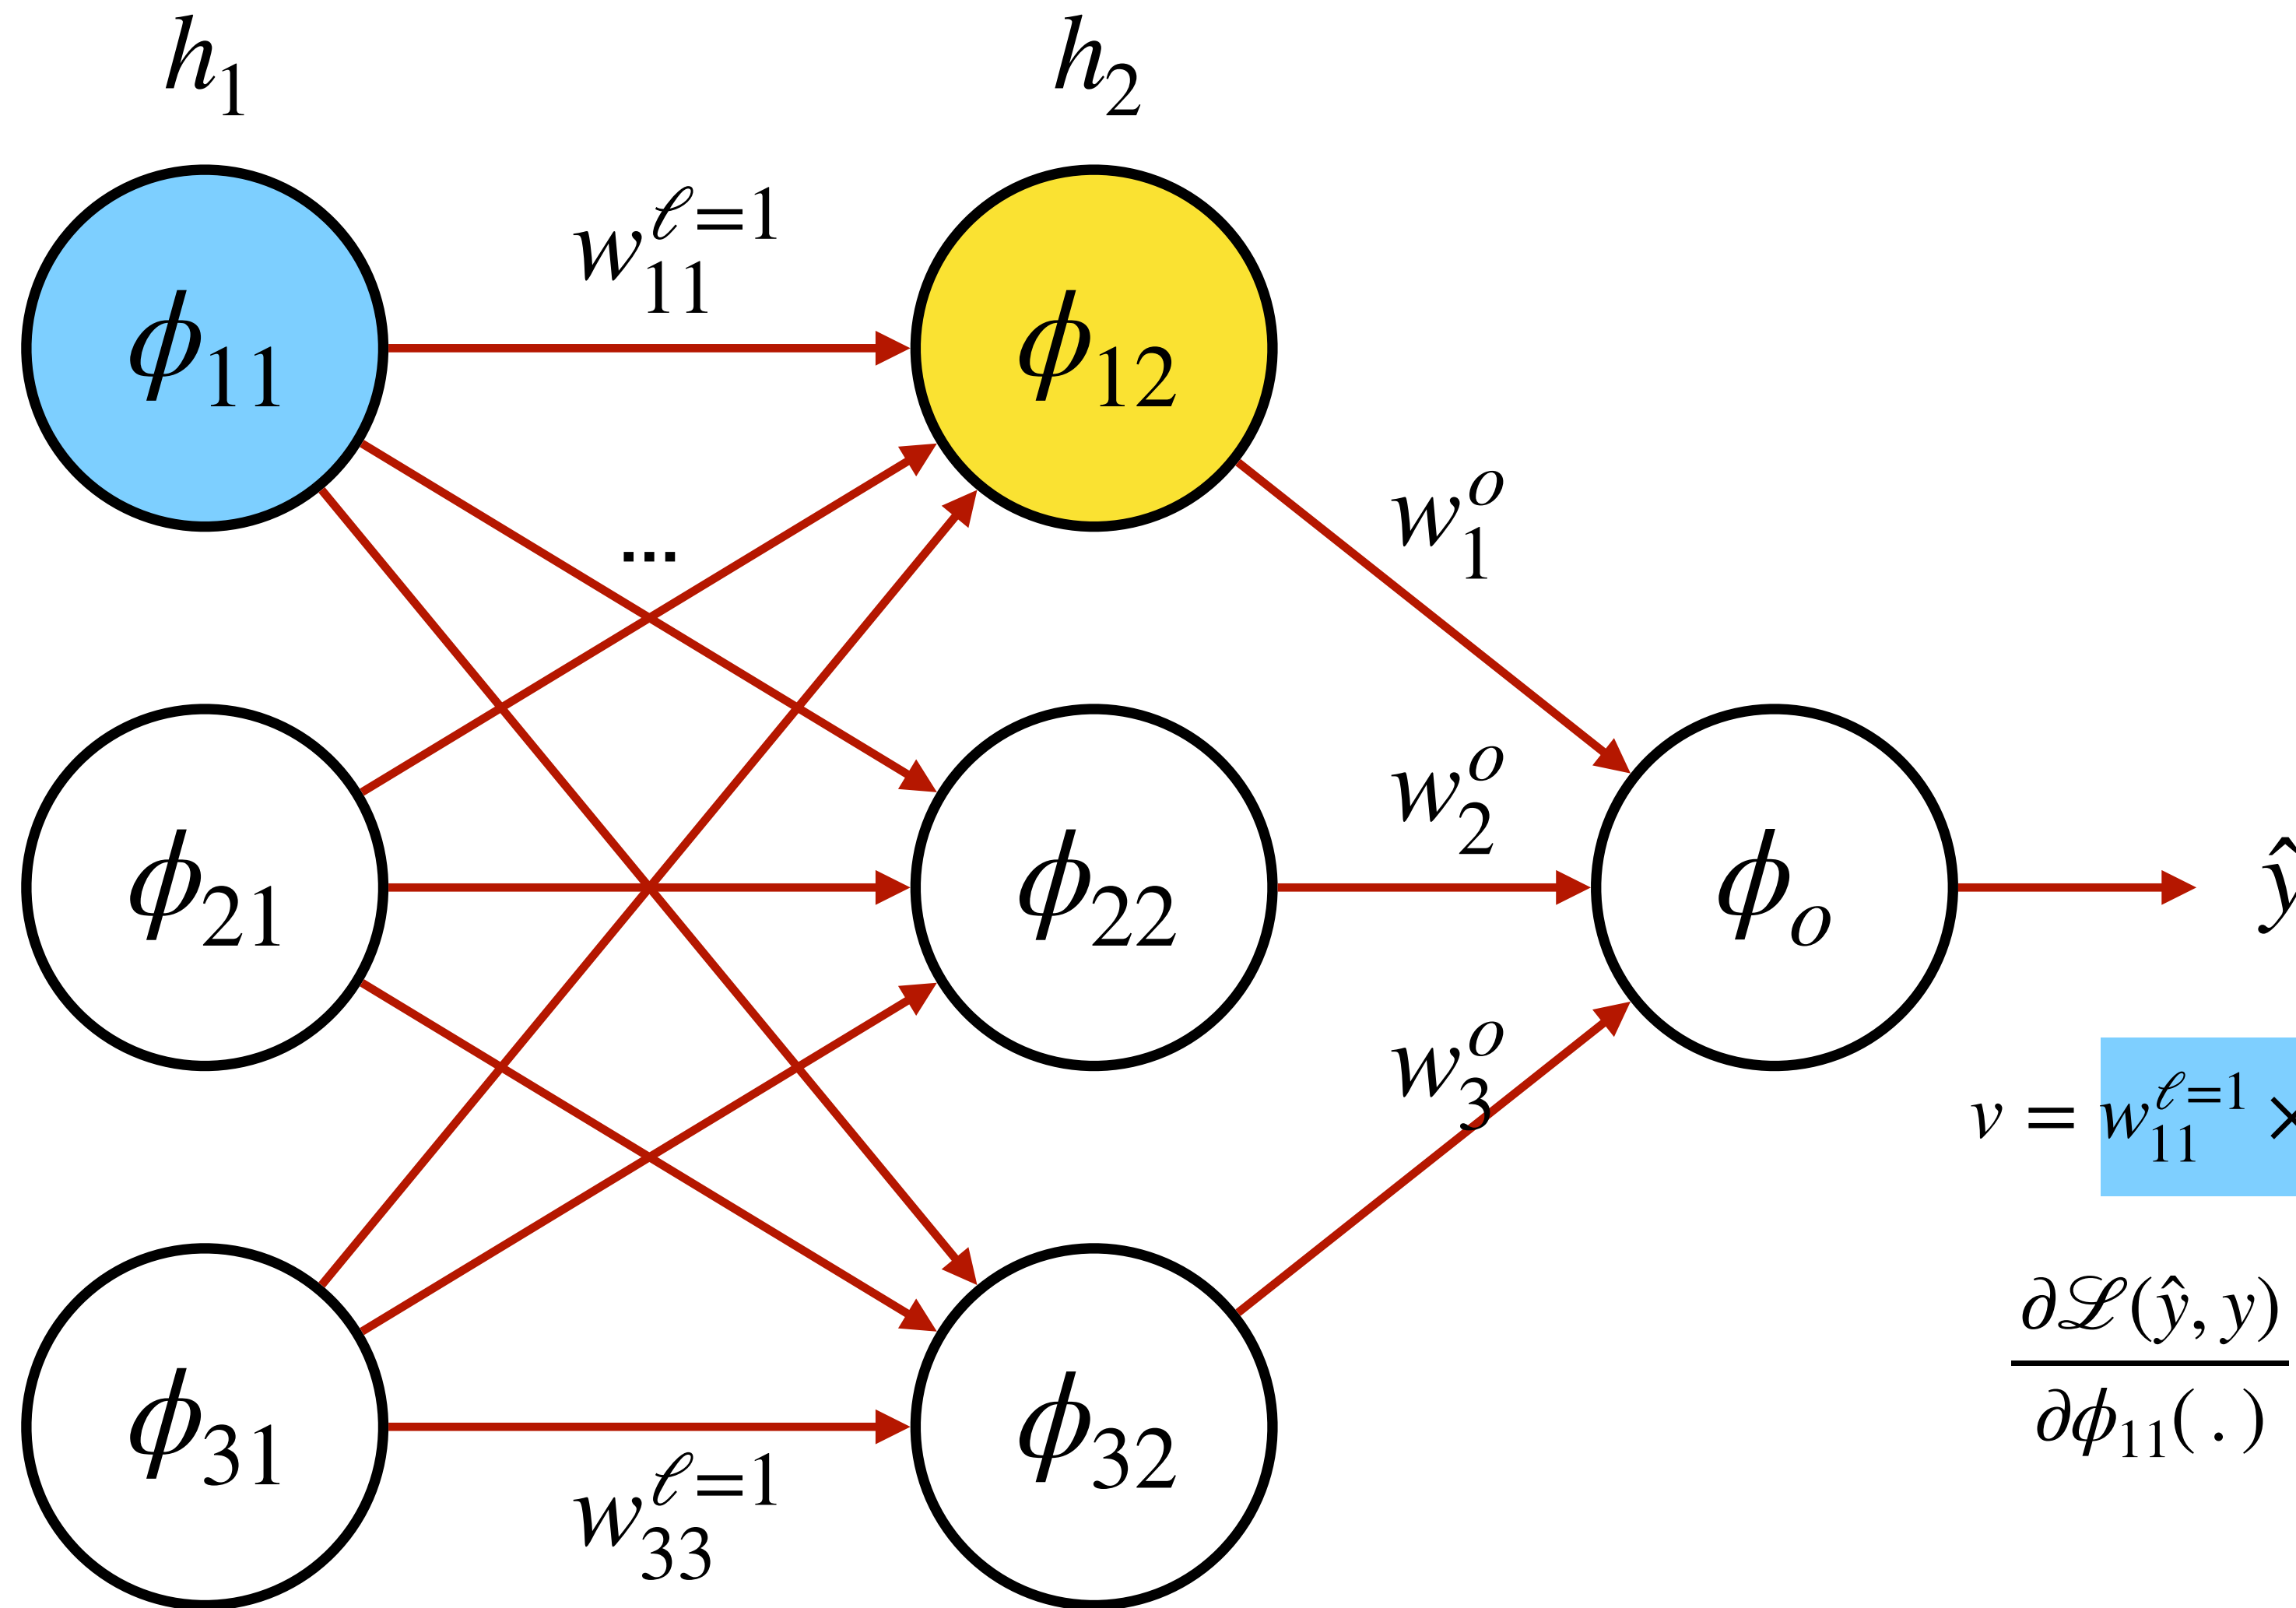$$= \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^0$$

Depends on ϕ_{12}

$$v = w_{11}^{\ell=1} \times \phi_{11}(.) + w_{21}^{\ell=1} \times \phi_{21}(.) + w_{31}^{\ell=1} \times \phi_{31}(.)$$

$$\frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \phi_{11}(.)} = \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(v)} \frac{\partial \phi_{12}(v)}{\partial v} \frac{\partial v}{\partial \phi_{11}(.)}$$

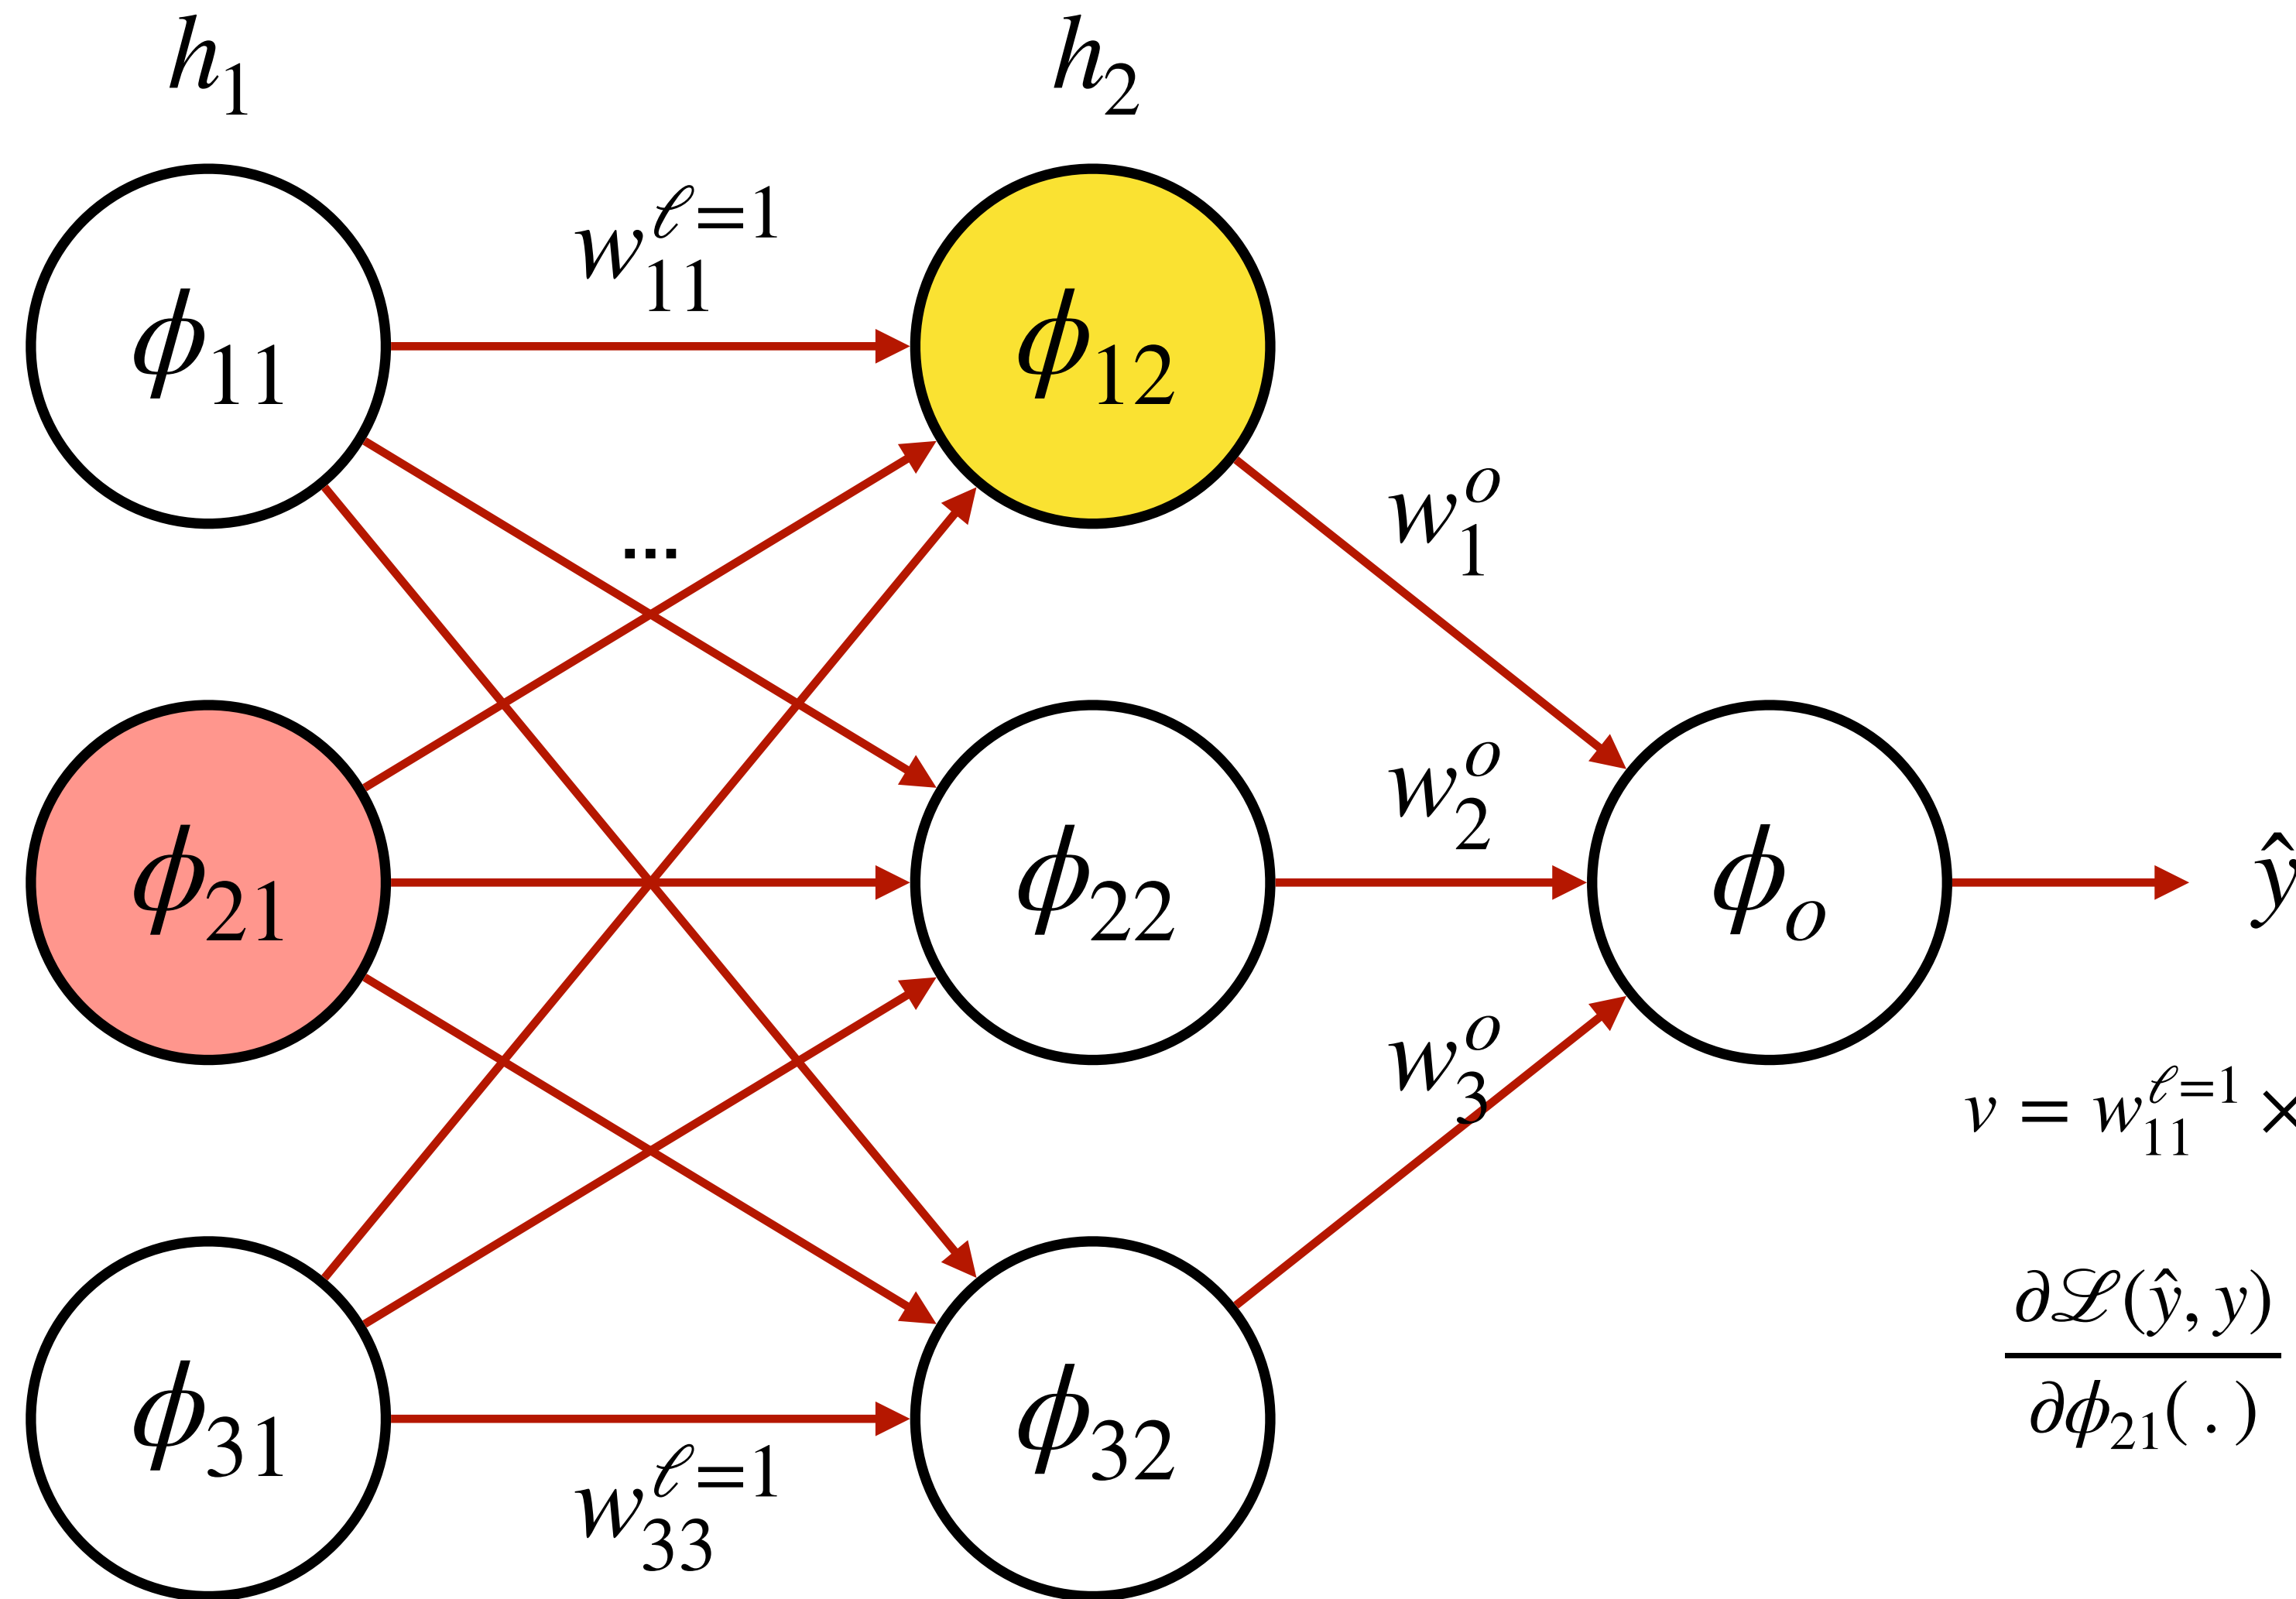$$= \frac{\partial \mathscr{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^0 \frac{\partial \phi_{12}(v)}{\partial v} w_{11}^{\ell=1}$$

Backpropagation Review: FFNs

Backpropagation Review: FFNs

Question

**How would we extend backpropagation
to a recurrent neural network?**

Recall

- RNN can be unrolled to a feedforward neural network

- Depth of feedforward neural network depends on length of the sequence

Backpropagation through Time

$$z_t = \sigma\left(W_{zh}h_t + b_z\right)$$

$$h_t = \sigma\left(W_{hx}x_t + W_{hh}h_{t-1} + b_h\right)$$

Backpropagation through Time

$$z_t = \sigma\big(W_{zh}h_t + b_z\big)$$

$$h_t = \sigma\big(W_{hx}x_t + W_{hh}h_{t-1} + b_h\big)$$

$$v = W_{zh}h_t + b_z \qquad\qquad z_t = \sigma(v)$$

$$u = W_{hx}x_t + W_{hh}h_{t-1} + b_h \qquad h_t = \sigma(u)$$

$$\frac{\partial z_t}{\partial h_t} = \frac{\partial \sigma(v)}{\partial v}\frac{\partial v}{\partial h_t} = \frac{\partial \sigma(v)}{\partial v}W_{zh}$$

Backpropagation through Time

$$z_t = \sigma\big(W_{zh}h_t + b_z\big)$$

$$h_t = \sigma\big(W_{hx}x_t + W_{hh}h_{t-1} + b_h\big)$$

$$v = W_{zh}h_t + b_z \qquad\qquad z_t = \sigma(v)$$

$$u = W_{hx}x_t + W_{hh}h_{t-1} + b_h \qquad h_t = \sigma(u)$$

$$\frac{\partial z_t}{\partial h_t} = \frac{\partial \sigma(v)}{\partial v}\frac{\partial v}{\partial h_t} = \frac{\partial \sigma(v)}{\partial v}W_{zh}$$

$$\frac{\partial h_t}{\partial h_{t-1}} = \frac{\partial \sigma(u)}{\partial u}\frac{\partial u}{\partial h_{t-1}} = \frac{\partial \sigma(u)}{\partial u}W_{hh}$$

Backpropagation through Time

$$z_t = \sigma\big(W_{zh}h_t + b_z\big)$$

$$h_t = \sigma\big(W_{hx}x_t + W_{hh}h_{t-1} + b_h\big)$$

$$v = W_{zh}h_t + b_z \qquad\qquad z_t = \sigma(v)$$

$$u = W_{hx}x_t + W_{hh}h_{t-1} + b_h \qquad h_t = \sigma(u)$$

$$\frac{\partial z_t}{\partial h_t} = \frac{\partial \sigma(v)}{\partial v}\frac{\partial v}{\partial h_t} = \frac{\partial \sigma(v)}{\partial v}W_{zh}$$

$$\frac{\partial h_t}{\partial h_{t-1}} = \frac{\partial \sigma(u)}{\partial u}\frac{\partial u}{\partial h_{t-1}} = \frac{\partial \sigma(u)}{\partial u}W_{hh}$$

$$\frac{\partial z_t}{\partial h_{t-1}} = \frac{\partial z_t}{\partial h_t}\frac{\partial h_t}{\partial h_{t-1}}$$

Backpropagation through Time

$$z_t = \sigma\left(W_{zh}h_t + b_z\right)$$

$$h_t = \sigma\left(W_{hx}x_t + W_{hh}h_{t-1} + b_h\right)$$

$$v = W_{zh}h_t + b_z \qquad z_t = \sigma(v)$$

$$u = W_{hx}x_t + W_{hh}h_{t-1} + b_h \qquad h_t = \sigma(u)$$

$$\frac{\partial z_t}{\partial h_t} = \frac{\partial \sigma(v)}{\partial v}\frac{\partial v}{\partial h_t} = \frac{\partial \sigma(v)}{\partial v}W_{zh}$$

$$\frac{\partial h_t}{\partial h_{t-1}} = \frac{\partial \sigma(u)}{\partial u}\frac{\partial u}{\partial h_{t-1}} = \frac{\partial \sigma(u)}{\partial u}W_{hh}$$

$$\frac{\partial z_t}{\partial h_{t-1}} = \frac{\partial z_t}{\partial h_t}\frac{\partial h_t}{\partial h_{t-1}} = \frac{\partial \sigma(v)}{\partial v}\frac{\partial v}{\partial h_t}\frac{\partial \sigma(u)}{\partial u}\frac{\partial u}{\partial h_{t-1}} = \frac{\partial \sigma(v)}{\partial v}W_{zh}\frac{\partial \sigma(u)}{\partial u}W_{hh}$$

Backpropagation through Time
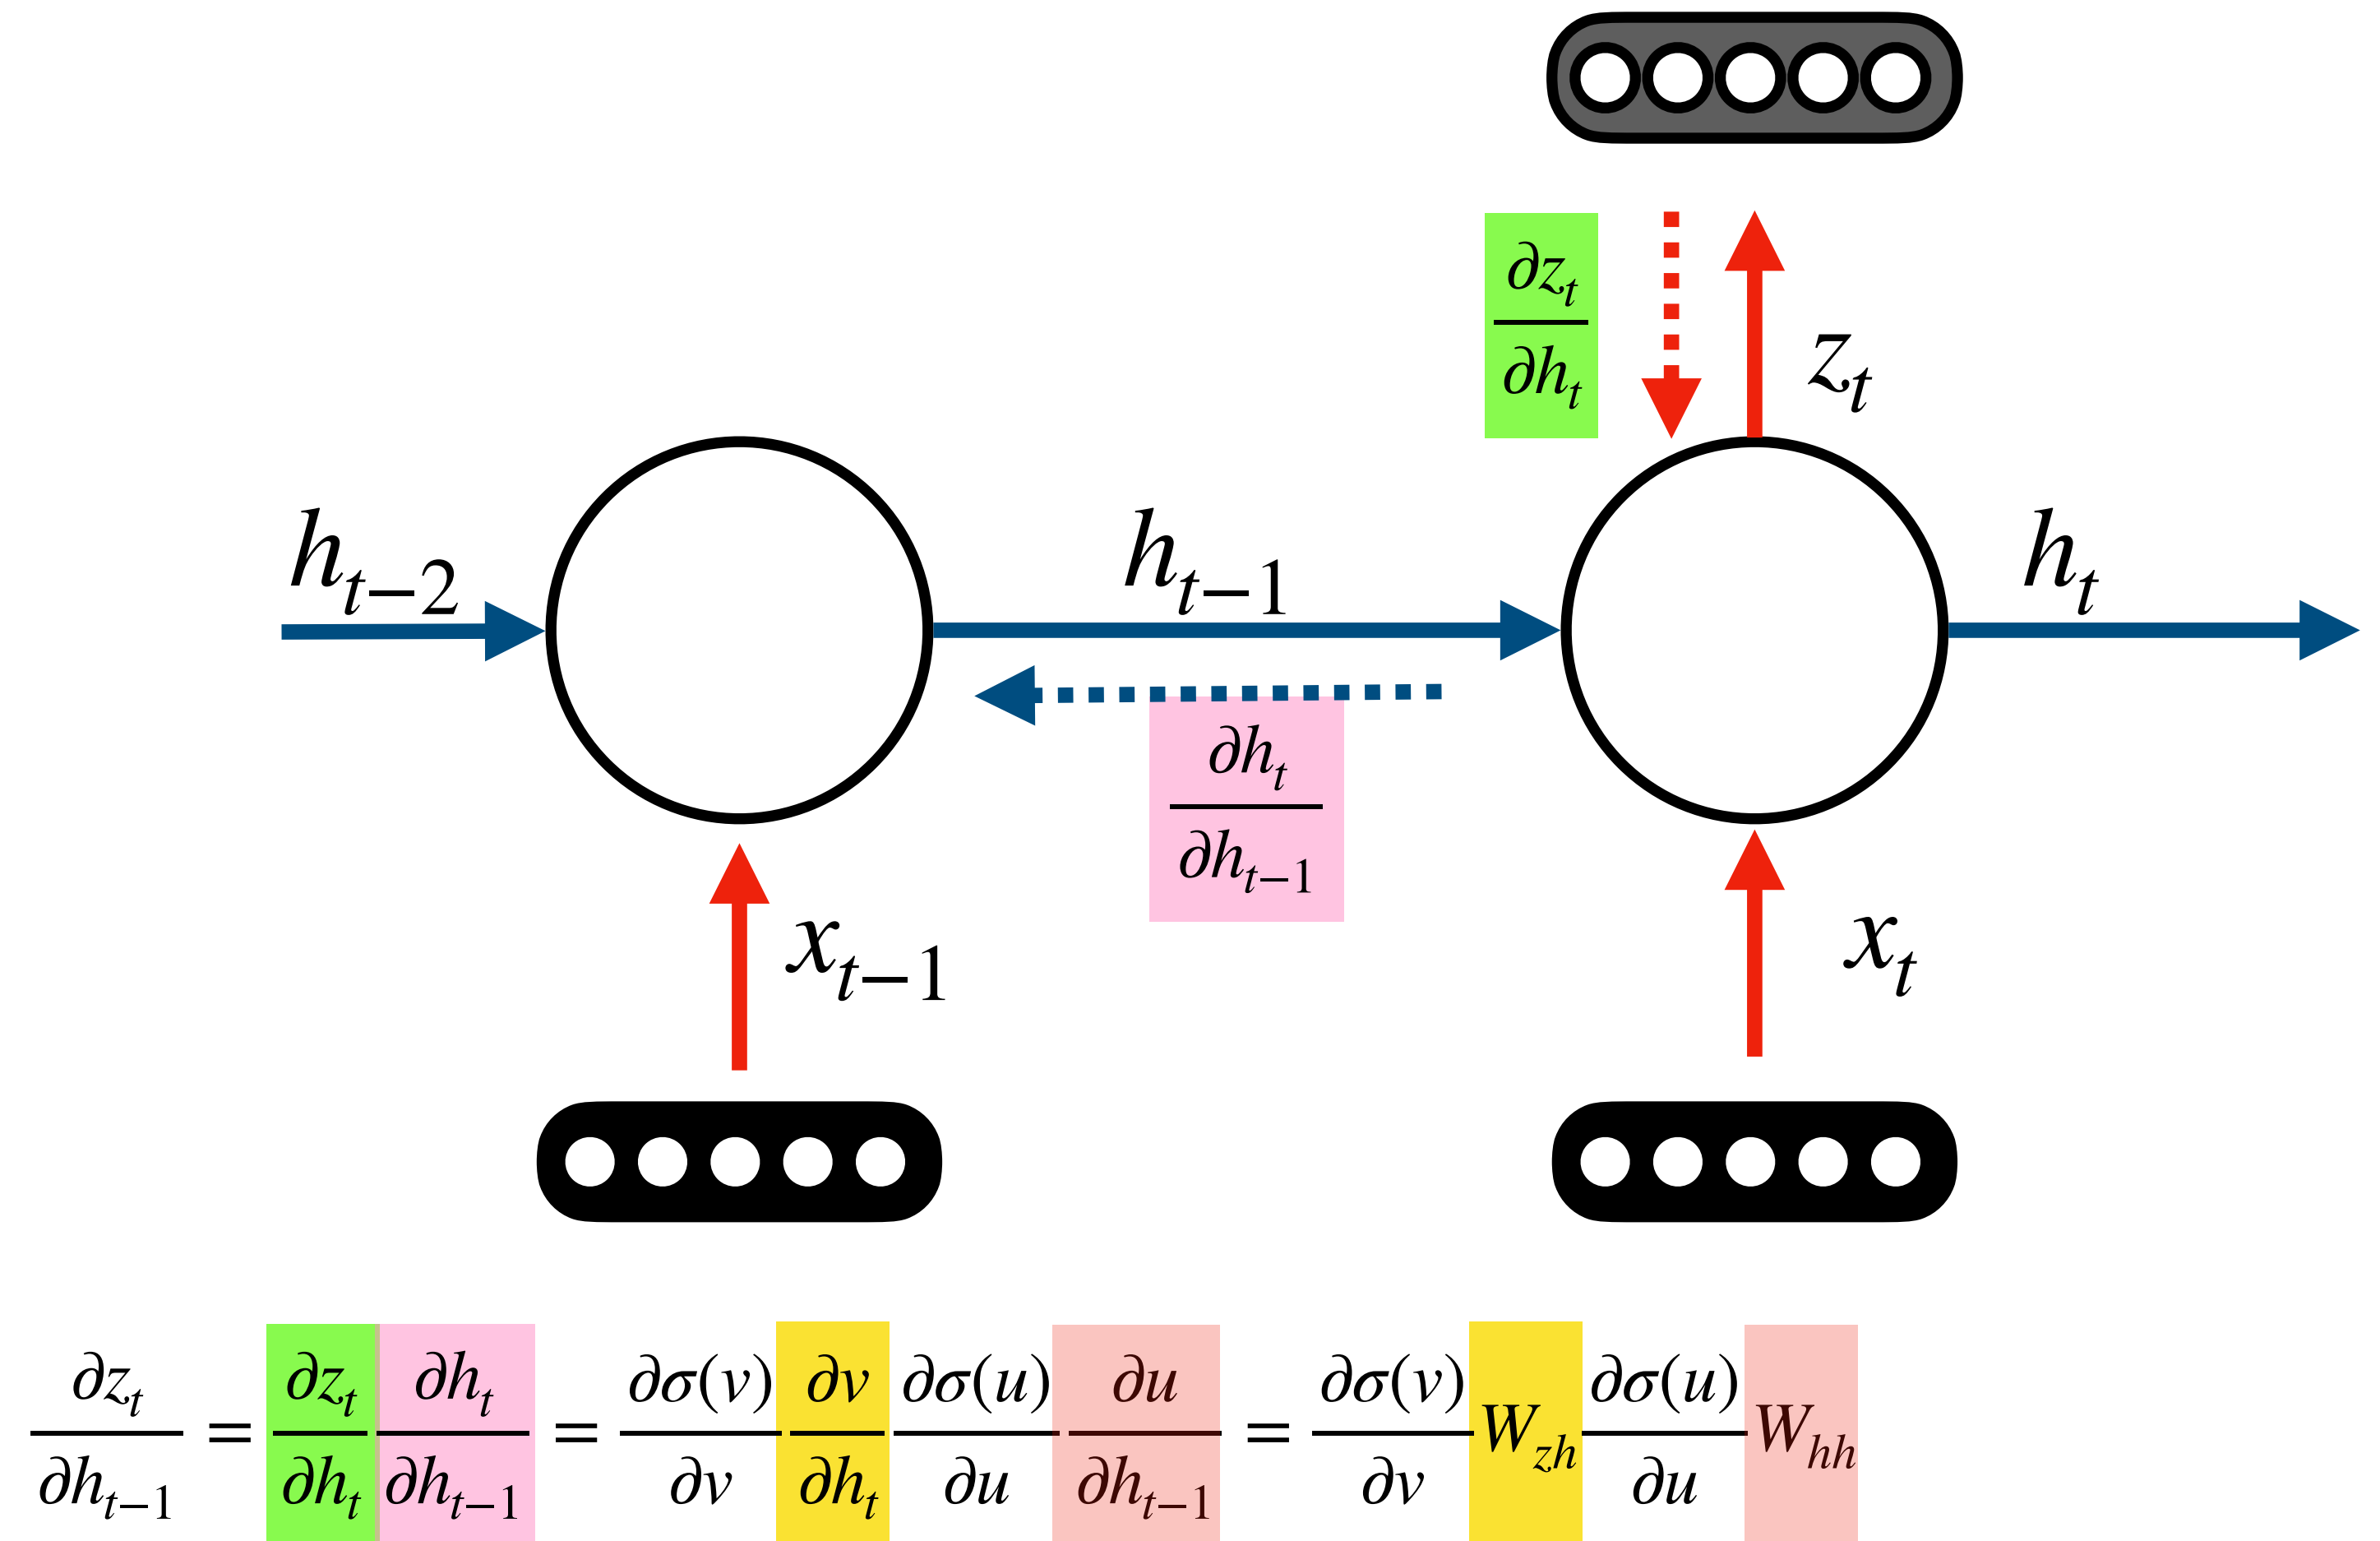
$$z_t = \sigma(W_{zh}h_t + b_z)$$

$$h_t = \sigma(W_{hx}x_t + W_{hh}h_{t-1} + b_h)$$

$$v_t = W_{zh}h_t + b_z \qquad z_t = \sigma(v_t)$$

$$u_t = W_{hx}x_t + W_{hh}h_{t-1} + b_h \qquad h_t = \sigma(u_t)$$

$$\frac{\partial z_t}{\partial h_t} = \frac{\partial \sigma(v_t)}{\partial v_t}\frac{\partial v_t}{\partial h_t} = \frac{\partial \sigma(v_t)}{\partial v_t}W_{zh}$$

$$\frac{\partial h_t}{\partial h_{t-1}} = \frac{\partial \sigma(u_t)}{\partial u_t}\frac{\partial u_t}{\partial h_{t-1}} = \frac{\partial \sigma(u_t)}{\partial u_t}W_{hh}$$

$$\frac{\partial h_{t-1}}{\partial h_{t-2}} = \frac{\partial \sigma(u_{t-1})}{\partial u_{t-1}}\frac{\partial u_{t-1}}{\partial h_{t-2}} = \frac{\partial \sigma(u_{t-1})}{\partial u_{t-1}}W_{hh}$$

$$\frac{\partial z_t}{\partial h_{t-1}} = \frac{\partial z_t}{\partial h_t}\frac{\partial h_t}{\partial h_{t-1}}\frac{\partial h_{t-1}}{\partial h_{t-2}} = \frac{\partial \sigma(v_t)}{\partial v_t}W_{zh}\frac{\partial \sigma(u_t)}{\partial u_t}W_{hh}\frac{\partial \sigma(u_{t-1})}{\partial u_{t-1}}W_{hh}$$

Backpropagation through Time

$$z_t = \sigma\left(W_{zh}h_t + b_z\right)$$

$$h_t = \sigma\left(W_{hx}x_t + W_{hh}h_{t-1} + b_h\right)$$


$$v_t = W_{zh}h_t + b_z \qquad\qquad z_t = \sigma(v_t)$$

$$u_t = W_{hx}x_t + W_{hh}h_{t-1} + b_h \qquad h_t = \sigma(u_t)$$

$$\frac{\partial z_t}{\partial h_t} = \frac{\partial \sigma(v_t)}{\partial v_t}\frac{\partial v_t}{\partial h_t} = \frac{\partial \sigma(v_t)}{\partial v_t}W_{zh}$$

$$\frac{\partial h_t}{\partial h_{t-1}} = \frac{\partial \sigma(u_t)}{\partial u_t}\frac{\partial u_t}{\partial h_{t-1}} = \frac{\partial \sigma(u_t)}{\partial u_t}W_{hh}$$

$$\frac{\partial h_{t-1}}{\partial h_{t-2}} = \frac{\partial \sigma(u_{t-1})}{\partial u_{t-1}}\frac{\partial u_{t-1}}{\partial h_{t-2}} = \frac{\partial \sigma(u_{t-1})}{\partial u_{t-1}}W_{hh}$$

$$\frac{\partial z_t}{\partial h_{t-1}} = \frac{\partial z_t}{\partial h_t}\frac{\partial h_t}{\partial h_{t-1}}\frac{\partial h_{t-1}}{\partial h_{t-2}} = \frac{\partial \sigma(v_t)}{\partial v_t}W_{zh}\frac{\partial \sigma(u_t)}{\partial u_t}W_{hh}\frac{\partial \sigma(u_{t-1})}{\partial u_{t-1}}W_{hh}$$

**Note that these are
actually the same matrix**

Backpropagation through time

Gradient flow

Output flow

mouse

$\frac{\partial h_3}{\partial h_2}$ $\frac{\partial h_4}{\partial h_3}$ $\frac{\partial h_5}{\partial h_4}$ $\frac{\partial h_6}{\partial h_5}$ $\frac{\partial h_7}{\partial h_6}$ z_7

h_2

h_3 h_4 h_5 h_6

$\frac{\partial h_3}{\partial x_3}$ x_3 $\frac{\partial h_4}{\partial x_4}$ x_4 $\frac{\partial h_5}{\partial x_5}$ x_5 $\frac{\partial h_6}{\partial x_6}$ x_6 $\frac{\partial h_7}{\partial x_7}$ x_7

cat   *fanatically*   *chased*   *the*   *elusive*
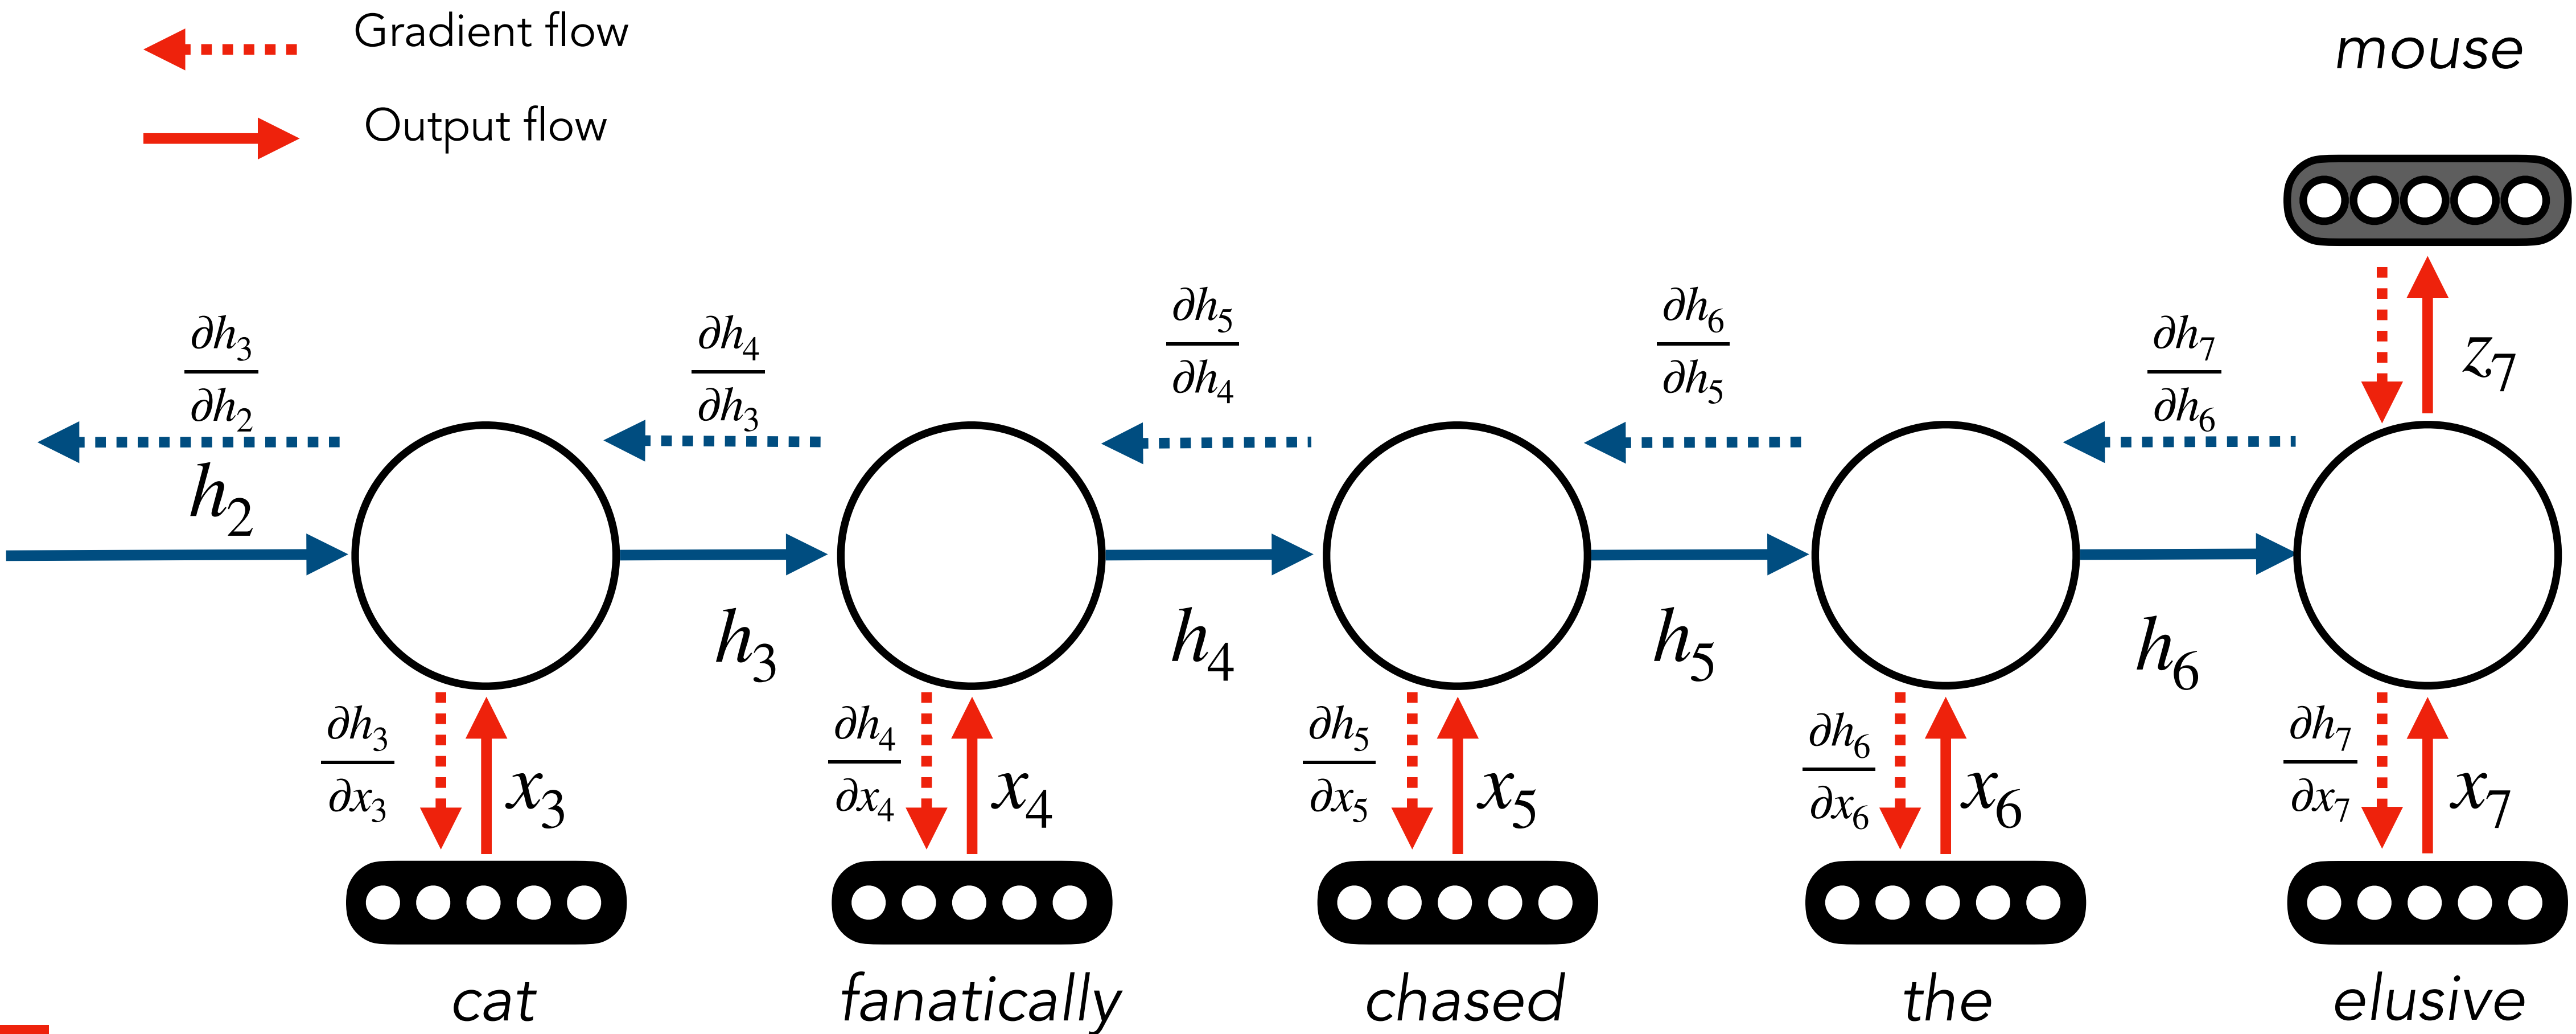
RNNs In Practice

- Computing gradients by hand is hard!

  - Though potentially a good way to sanity check whether you network behaves the way you expect

- Most modern software packages for deep learning use automatic differentiation to compute gradients automatically from the forward pass

- Only need to define the forward pass of your model (much easier!)

- You'll use PyTorch in this class! You won't have to compute gradients by hand :)

Multiple Layers

In practice, typically use 3-4 layers of RNN cells

Bidirectionality

$[z_3^f; z_3^b]$ $[z_4^f; z_4^b]$ $[z_5^f; z_5^b]$ $[z_6^f; z_6^b]$ $[z_7^f; z_7^b]$

h_2^b h_3^b h_4^b h_5^b h_6^b h_7^b

h_2^f h_3^f h_4^f h_5^f h_6^f h_7^f

x_3 x_4 x_5 x_6 x_7

cat *fanatically* *chased* *the* *elusive*
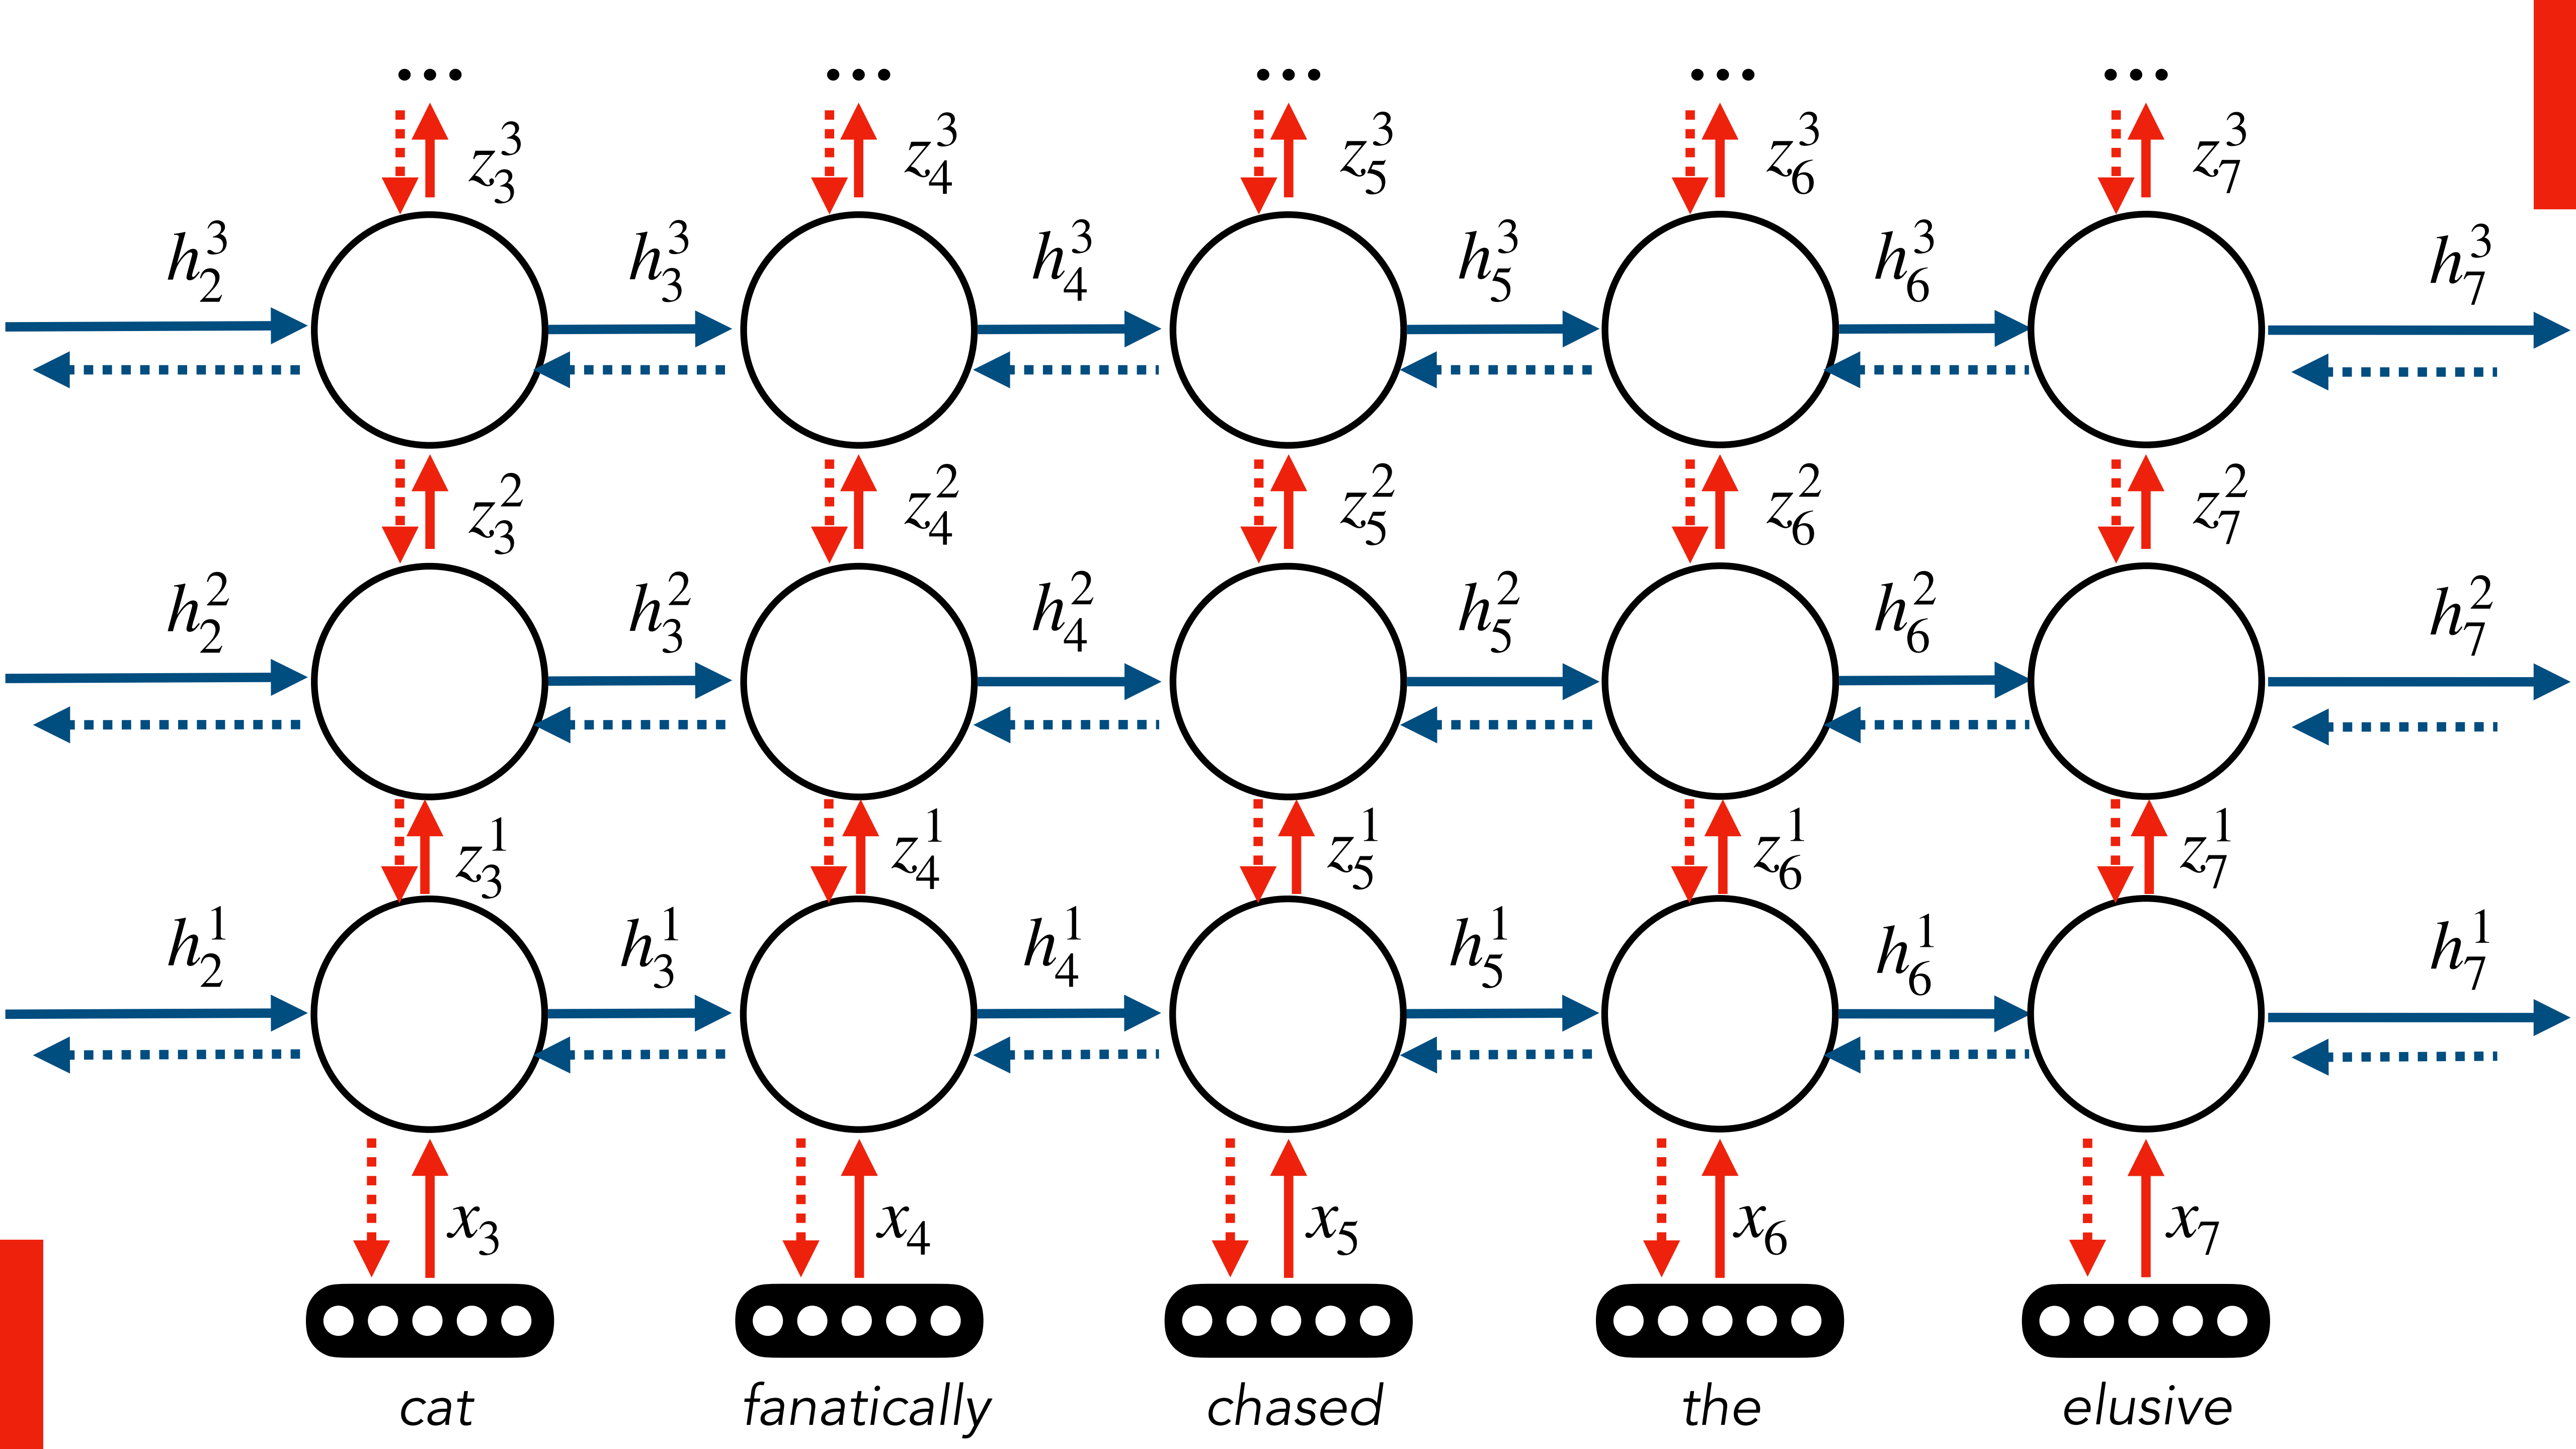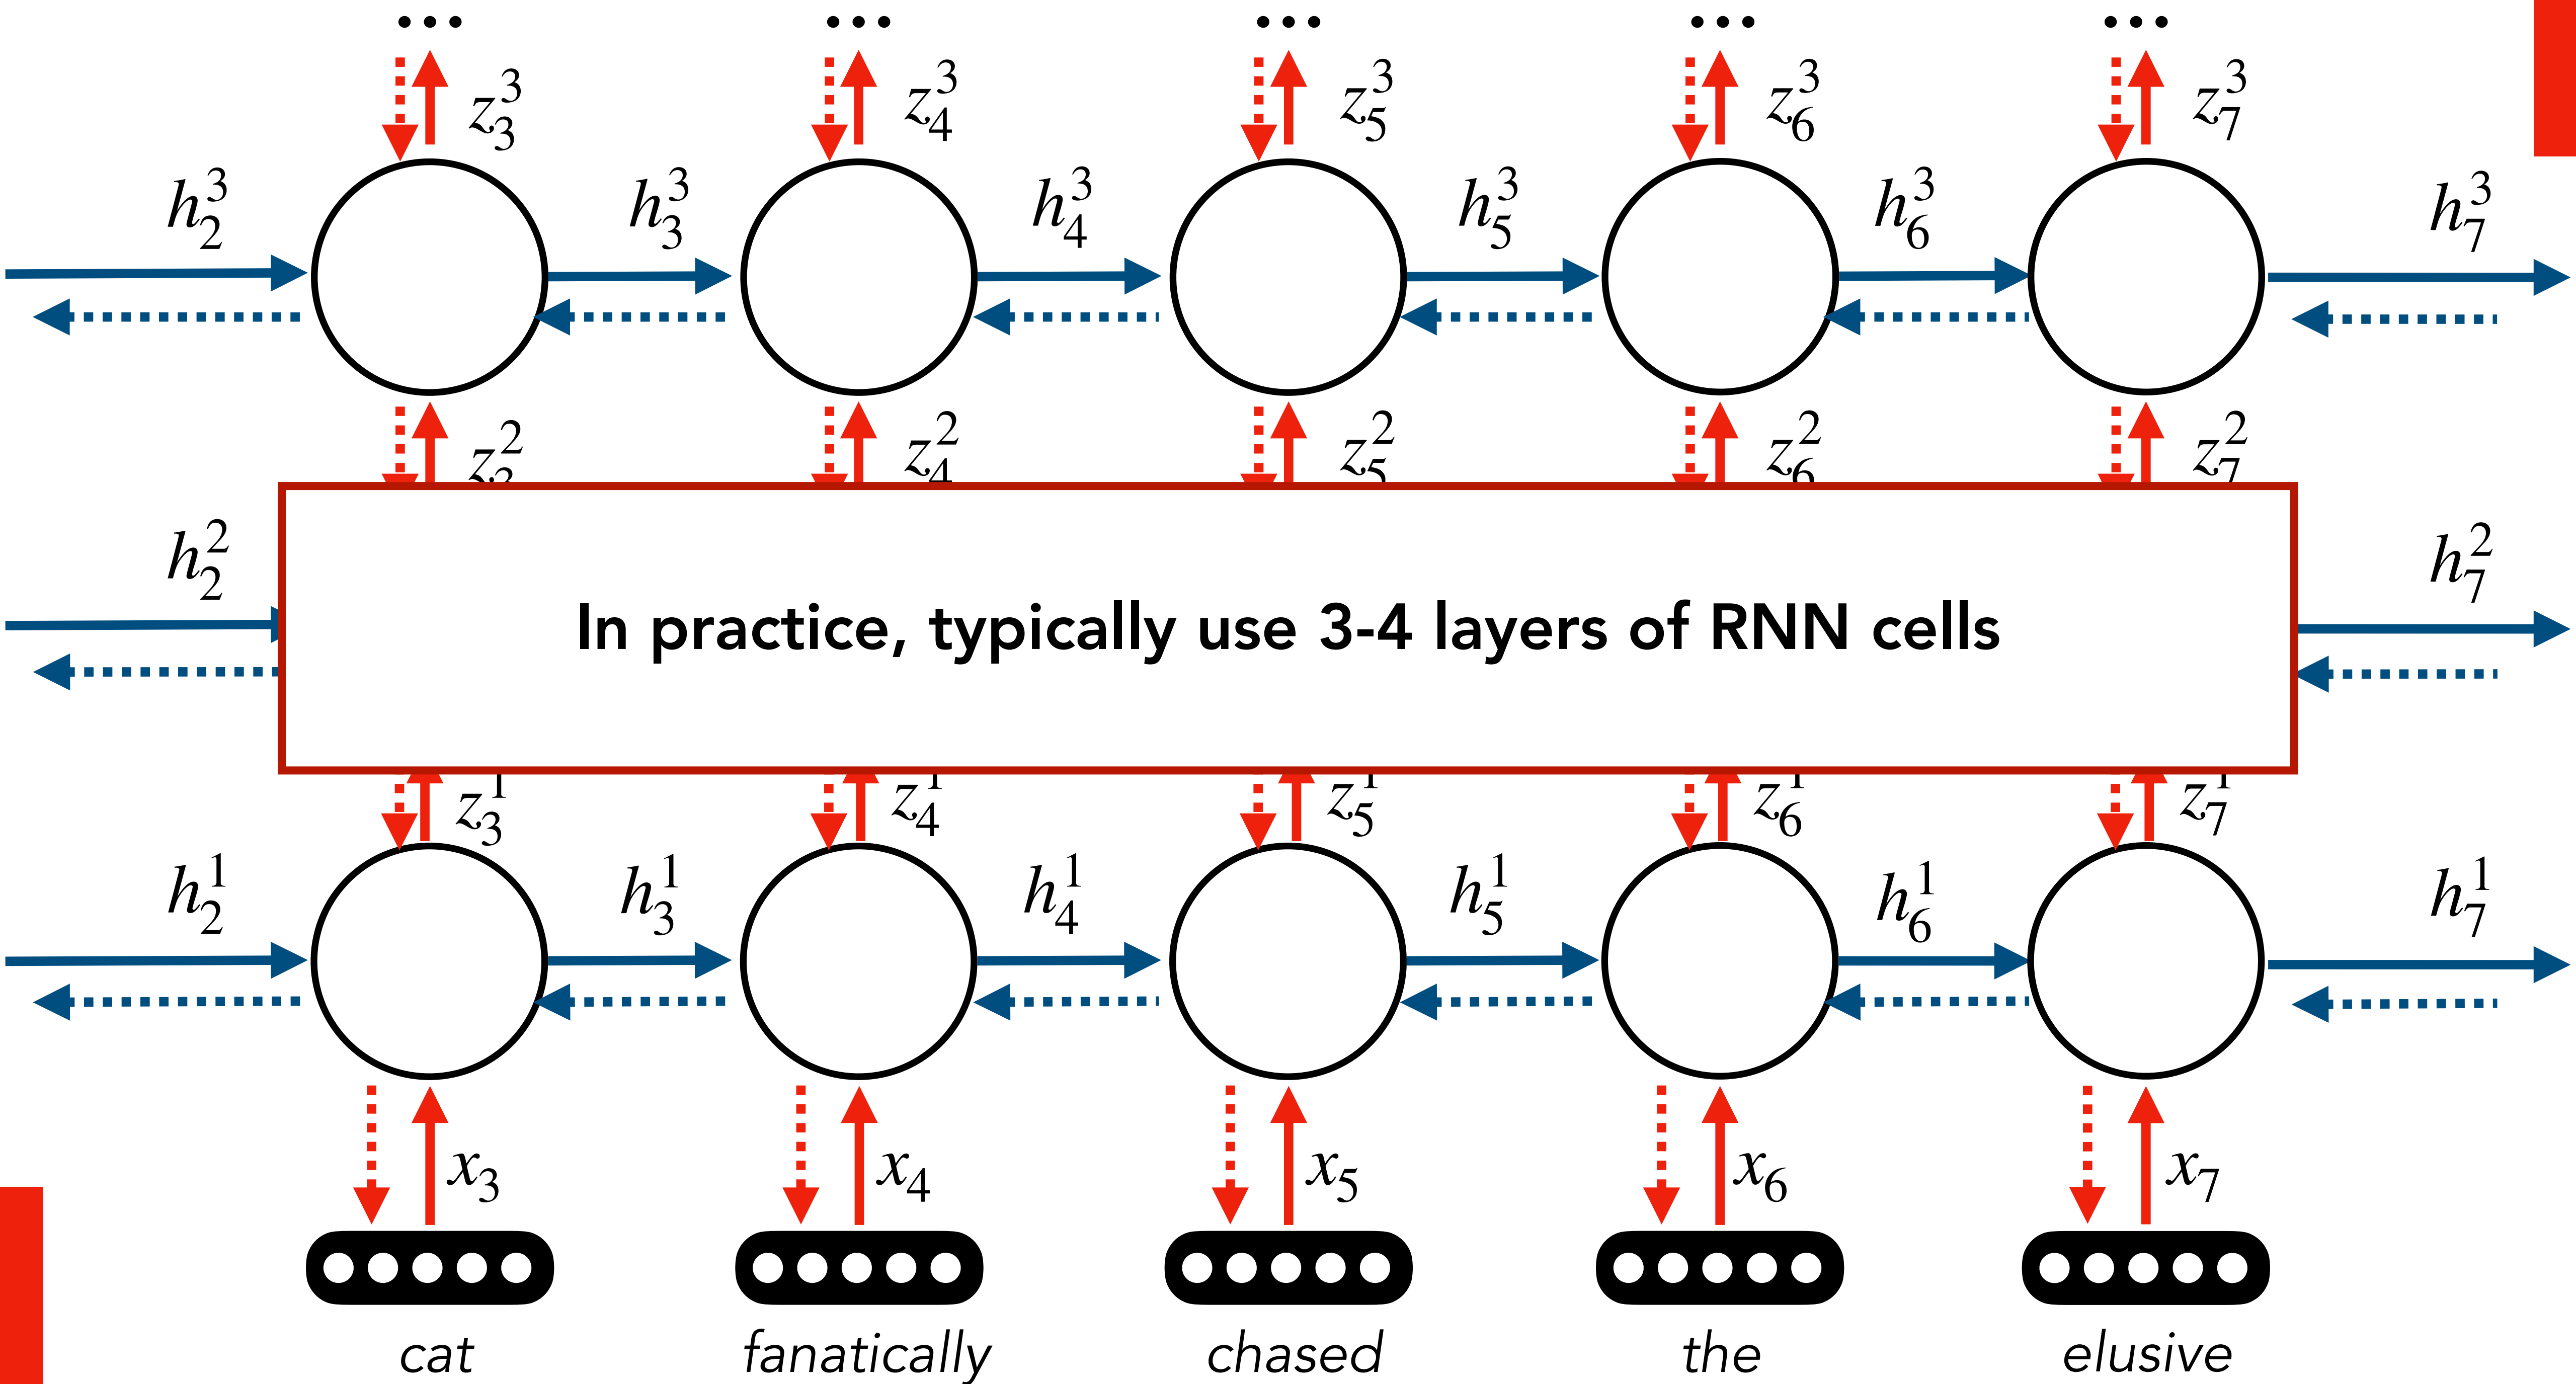
Bidirectionality

- **Concatenate** the output states for final representation at each step
  - Can also do mean / max

- Separate parameters for forward and backward RNNs

- If you can use the future text for the task, then you should use **bidirectionality**

Question

For which of the following task types
can we use a bidirectional RNN?

(a)   Classification
(b)   Sequence labelling
(c)   Text Generation

Recap

- Neural language models allow us to *share information* among similar sequences by learning neural representations that similarly represent them

- **Problem:** Fixed context language models can only process a limited window of the word history at a time

- **Solution:** recurrent neural networks can **theoretically** learn to model an **unbounded context length**

Issue with Recurrent Models

- Multiple steps of state overwriting makes it challenging to learn long-range dependencies.

  *They tuned, discussed for a moment, then struck up a lively **jig**. Everyone joined in, turning the courtyard into an even more chaotic scene, people now **dancing** in circles, **swinging** and **spinning** in circles, everyone making up their own **dance steps**. I felt my feet tapping, my body wanting to move. Aside from writing, I 've always loved **dancing** .*

- Nearby words should affect each other more than farther ones, but RNNs make it challenging to learn **<u>any</u>** long-range interactions

LAMBADA dataset, 2016

Backpropagation through time

Vanishing Gradients

$$z_t = \sigma\left(W_{zh}h_t + b_z\right)$$

$$h_t = \sigma\left(W_{hx}x_t + W_{hh}h_{t-1} + b_h\right)$$

$$v_t = W_{zh}h_t + b_z \qquad\qquad z_t = \sigma(v_t)$$

$$u_t = W_{hx}x_t + W_{hh}h_{t-1} + b_h \qquad h_t = \sigma(u_t)$$

$$\frac{\partial z_t}{\partial h_t} = \frac{\partial \sigma(v_t)}{\partial v_t}\frac{\partial v_t}{\partial h_t} = \frac{\partial \sigma(v_t)}{\partial v_t}W_{zh}$$

$$\frac{\partial h_t}{\partial h_{t-1}} = \frac{\partial \sigma(u_t)}{\partial u_t}\frac{\partial u_t}{\partial h_{t-1}} = \frac{\partial \sigma(u_t)}{\partial u_t}W_{hh}$$

$$\frac{\partial h_{t-1}}{\partial h_{t-2}} = \frac{\partial \sigma(u_{t-1})}{\partial u_{t-1}}\frac{\partial u_{t-1}}{\partial h_{t-2}} = \frac{\partial \sigma(u_{t-1})}{\partial u_{t-1}}W_{hh}$$

$$\frac{\partial z_t}{\partial h_{t-2}} = \frac{\partial z_t}{\partial h_t}\frac{\partial h_t}{\partial h_{t-1}}\frac{\partial h_{t-1}}{\partial h_{t-2}} = \frac{\partial \sigma(v_t)}{\partial v_t}W_{zh}\frac{\partial \sigma(u_t)}{\partial u_t}W_{hh}\frac{\partial \sigma(u_{t-1})}{\partial u_{t-1}}W_{hh}$$

Generalising this:

$$\frac{\partial h_t}{\partial h_{t-T}} = \prod_{i=t-T}^{i=t}\frac{\partial h_i}{\partial h_{i-1}} = \prod_{i=t-T}^{i=t}\frac{\partial \sigma(u_i)}{\partial u_i}W_{hh}$$

Vanishing Gradients

$$z_t = \sigma\left(W_{zh}h_t + b_z\right)$$

$$h_t = \sigma\left(W_{hx}x_t + W_{hh}h_{t-1} + b_h\right)$$

$$v_t = W_{zh}h_t + b_z \qquad\qquad z_t = \sigma(v_t)$$

$$u_t = W_{hx}x_t + W_{hh}h_{t-1} + b_h \qquad h_t = \sigma(u_t)$$

$$\frac{\partial z_t}{\partial h_t} = \frac{\partial\sigma(v_t)}{\partial v_t}\frac{\partial v_t}{\partial h_t} = \frac{\partial\sigma(v_t)}{\partial v_t}W_{zh}$$

$$\frac{\partial h_t}{\partial h_{t-1}} = \frac{\partial\sigma(u_t)}{\partial u_t}\frac{\partial u_t}{\partial h_{t-1}} = \frac{\partial\sigma(u_t)}{\partial u_t}W_{hh}$$

$$\frac{\partial h_{t-1}}{\partial h_{t-2}} = \frac{\partial\sigma(u_{t-1})}{\partial u_{t-1}}\frac{\partial u_{t-1}}{\partial h_{t-2}} = \frac{\partial\sigma(u_{t-1})}{\partial u_{t-1}}W_{hh}$$

$$\frac{\partial z_t}{\partial h_{t-2}} = \frac{\partial z_t}{\partial h_t}\frac{\partial h_t}{\partial h_{t-1}}\frac{\partial h_{t-1}}{\partial h_{t-2}} = \frac{\partial\sigma(v_t)}{\partial v_t}W_{zh}\frac{\partial\sigma(u_t)}{\partial u_t}W_{hh}\frac{\partial\sigma(u_{t-1})}{\partial u_{t-1}}W_{hh}$$

Generalising this:

$$\frac{\partial h_t}{\partial h_{t-T}} = \prod_{i=t-T}^{i=t}\frac{\partial h_i}{\partial h_{i-1}} = \prod_{i=t-T}^{i=t}\frac{\partial\sigma(u_i)}{\partial u_i}W_{hh}$$

< 1 for many activation fxns

Typically small (Regularisation)

Vanishing Gradients

- **Learning Problem:** Long unrolled networks will crush gradients that backpropagate to earlier time steps

$$h_t = \sigma \left(W_{hx} x_t + W_{hh} h_{t-1} + b_h \right)$$

$$u = W_{hx} x_t + W_{hh} h_{t-1} + b_h$$

$$\frac{\partial h_t}{\partial h_{t-1}} = \frac{\partial \sigma(u)}{\partial u} \frac{\partial u}{\partial h_{t-1}} = W_{hh} \frac{\partial \sigma(u)}{\partial u}$$

Vanishing Gradients

Gradient flow

Output flow

$$\frac{\partial h_t}{\partial h_{t-1}} = \frac{\partial \sigma(u)}{\partial u} \frac{\partial u}{\partial h_{t-1}} = W_{hh} \frac{\partial \sigma(u)}{\partial u}$$

$\frac{\partial h_3}{\partial h_2}$ $\frac{\partial h_4}{\partial h_3}$ $\frac{\partial h_5}{\partial h_4}$ $\frac{\partial h_6}{\partial h_5}$ $\frac{\partial h_7}{\partial h_6}$

z_7

h_2

h_3 h_4 h_5 h_6

x_3 x_4 x_5 x_6 x_7

cat fanatically chased the elusive

Vanishing Gradients

Gradient flow

Output flow

$$\frac{\partial h_t}{\partial h_{t-1}} = \frac{\partial \sigma(u)}{\partial u}\frac{\partial u}{\partial h_{t-1}} = W_{hh}\frac{\partial \sigma(u)}{\partial u}$$

$\frac{\partial h_3}{\partial h_2}$

h_2

Problem in many recurrent neural networks,

**Especially pronounced in Elman networks
(Vanilla RNNs) due to the sigmoid activation**

z_7

x_3 x_4 x_5 x_6 x_7

cat   *fanatically*   *chased*   *the*   *elusive*

Question

How could we fix this vanishing gradient problem?

$$\frac{\partial h_t}{\partial h_{t-T}} = \prod_{i=t-T}^{i=t} \frac{\partial h_i}{\partial h_{i-1}} = \prod_{i=t-T}^{i=t} \frac{\partial \sigma(u_i)}{\partial u_i} W_{hh}$$

Typically less than one

Typically small (Regularisation)

Gated Recurrent Neural Networks

- Use gates to avoid dampening gradient signal every time step

$$h_t = \sigma\left(W_{hx}x_t + W_{hh}h_{t-1} + b_h\right) \qquad\qquad h_t = h_{t-1} \odot \mathbf{f} + \mathbf{func}(x_t)$$

<div align="center">

Elman Network   **Gated Network Abstraction**

</div>

- Gate value $\mathbf{f}$ computes how much information from previous hidden state moves to the next time step —> $0 < \mathbf{f} < 1$

- Because h_{t-1} is no longer inside the activation function, it is not automatically constrained, reducing vanishing gradients!

Long Short Term Memory (LSTM)

Gates:

$$f_t = \sigma\big(W_{fx}x_t + W_{fh}h_{t-1} + b_f\big)$$
$$i_t = \sigma\big(W_{ix}x_t + W_{ih}h_{t-1} + b_i\big)$$
$$o_t = \sigma\big(W_{ox}x_t + W_{oh}h_{t-1} + b_o\big)$$

$$\tilde{c}_t = \phi\big(W_{cx}x_t + W_{ch}h_{t-1} + b_c\big)$$

$$c_t = i_t \times \tilde{c}_t + f_t \times c_{t-1}$$

$$h_t = o_t \times \phi(c_t)$$

(Hochreiter and Schmidhuber, 1997)

Cell State

$$\tilde{c}_t = \phi\big(W_{cx}x_t + W_{ch}h_{t-1} + b_c\big)$$

$$c_t = i_t \times \tilde{c}_t + f_t \times c_{t-1}$$

- Hidden state h_{t-1} is now short-term memory

- Cell state c_t tracks longer-term dependencies

48

What does the cell state track?

- Can visualise the activations of cell state (i.e., dimensions of **c**) and find semantic behaviour!

- Stack Overflow example:

```
#ifdef CONFIG_AUDITSYSCALL
static inline int audit_match_class_bits(int class, u32 *mask)
{
  int i;
  if (classes[class]) {
    for (i = 0; i < AUDIT_BITMASK_SIZE; i++)
      if (mask[i] & classes[class][i])
        return 0;
  }
  return 1;
}
```

What does the cell state track?

- Can visualise the activations of cell state (i.e., dimensions of **c**) and find semantic behaviour!

- Stack Overflow example: **track indentation**

```
#ifdef CONFIG_AUDITSYSCALL
static inline int audit_match_class_bits(int class, u32 *mask)
{
    int i;
    if (classes[class]) {
        for (i = 0; i < AUDIT_BITMASK_SIZE; i++)
            if (mask[i] & classes[class][i])
                return 0;
    }
    return 1;
}
```

Karpathy et al. (2015)

What does the cell state track?

- Can visualise the activations of cell state (i.e., dimensions of **c**) and find semantic behaviour!

- Stack Overflow example: track indentation

```
#ifdef CONFIG_AUDITSYSCALL
static inline int audit_match_class_bits(int class, u32 *mask)
{
  int i;
  if (classes[class]) {
    for (i = 0; i < AUDIT_BITMASK_SIZE; i++)
      if (mask[i] & classes[class][i])
        return 0;
  }
  return 1;
}
```

- War and Peace:

```
"You mean to imply that I have nothing to eat out of.... On the
contrary, I can supply you with everything even if you want to give
dinner parties," warmly replied Chichagov, who tried by every word he
spoke to prove his own rectitude and therefore imagined Kutuzov to be
animated by the same desire.

Kutuzov, shrugging his shoulders, replied with his subtle penetrating
smile: "I meant merely to say what I said."
```

Karpathy et al. (2015)

What does the cell state track?

- Can visualise the activations of cell state (i.e., dimensions of **c**) and find semantic behaviour!

- Stack Overflow example: track indentation

```
#ifdef CONFIG_AUDITSYSCALL
static inline int audit_match_class_bits(int class, u32 *mask)
{
  int i;
  if (classes[class]) {
    for (i = 0; i < AUDIT_BITMASK_SIZE; i++)
      if (mask[i] & classes[class][i])
        return 0;
  }
  return 1;
}
```

- War and Peace: **are we in a quote or not?**

```
"You mean to imply that I have nothing to eat out of.... On the
contrary, I can supply you with everything even if you want to give
dinner parties," warmly replied Chichagov, who tried by every word he
spoke to prove his own rectitude and therefore imagined Kutuzov to be
animated by the same desire.

Kutuzov, shrugging his shoulders, replied with his subtle penetrating
smile: "I meant merely to say what I said."
```

Karpathy et al. (2015)

Long Short Term Memory (LSTM)

Gates:

$$f_t = \sigma\big(W_{fx}x_t + W_{fh}h_{t-1} + b_f\big)$$
$$i_t = \sigma\big(W_{ix}x_t + W_{ih}h_{t-1} + b_i\big)$$
$$o_t = \sigma\big(W_{ox}x_t + W_{oh}h_{t-1} + b_o\big)$$

$$\tilde{c}_t = \phi\big(W_{cx}x_t + W_{ch}h_{t-1} + b_c\big)$$

$$c_t = i_t \times \tilde{c}_t + f_t \times c_{t-1}$$

$$h_t = o_t \times \phi(c_t)$$

Forget Gate

I went to **the** **lecture**

$$f_t = \sigma\big(W_{fx}x_t + W_{fh}h_{t-1} + b_f\big)$$

$$\tilde{c}_t = \phi\big(W_{cx}x_t + W_{ch}h_{t-1} + b_c\big)$$

$$c_t = i_t \times \tilde{c}_t + f_t \times c_{t-1}$$

- Forget gate controls how much memory is forgotten

  - 1 -> remember the past

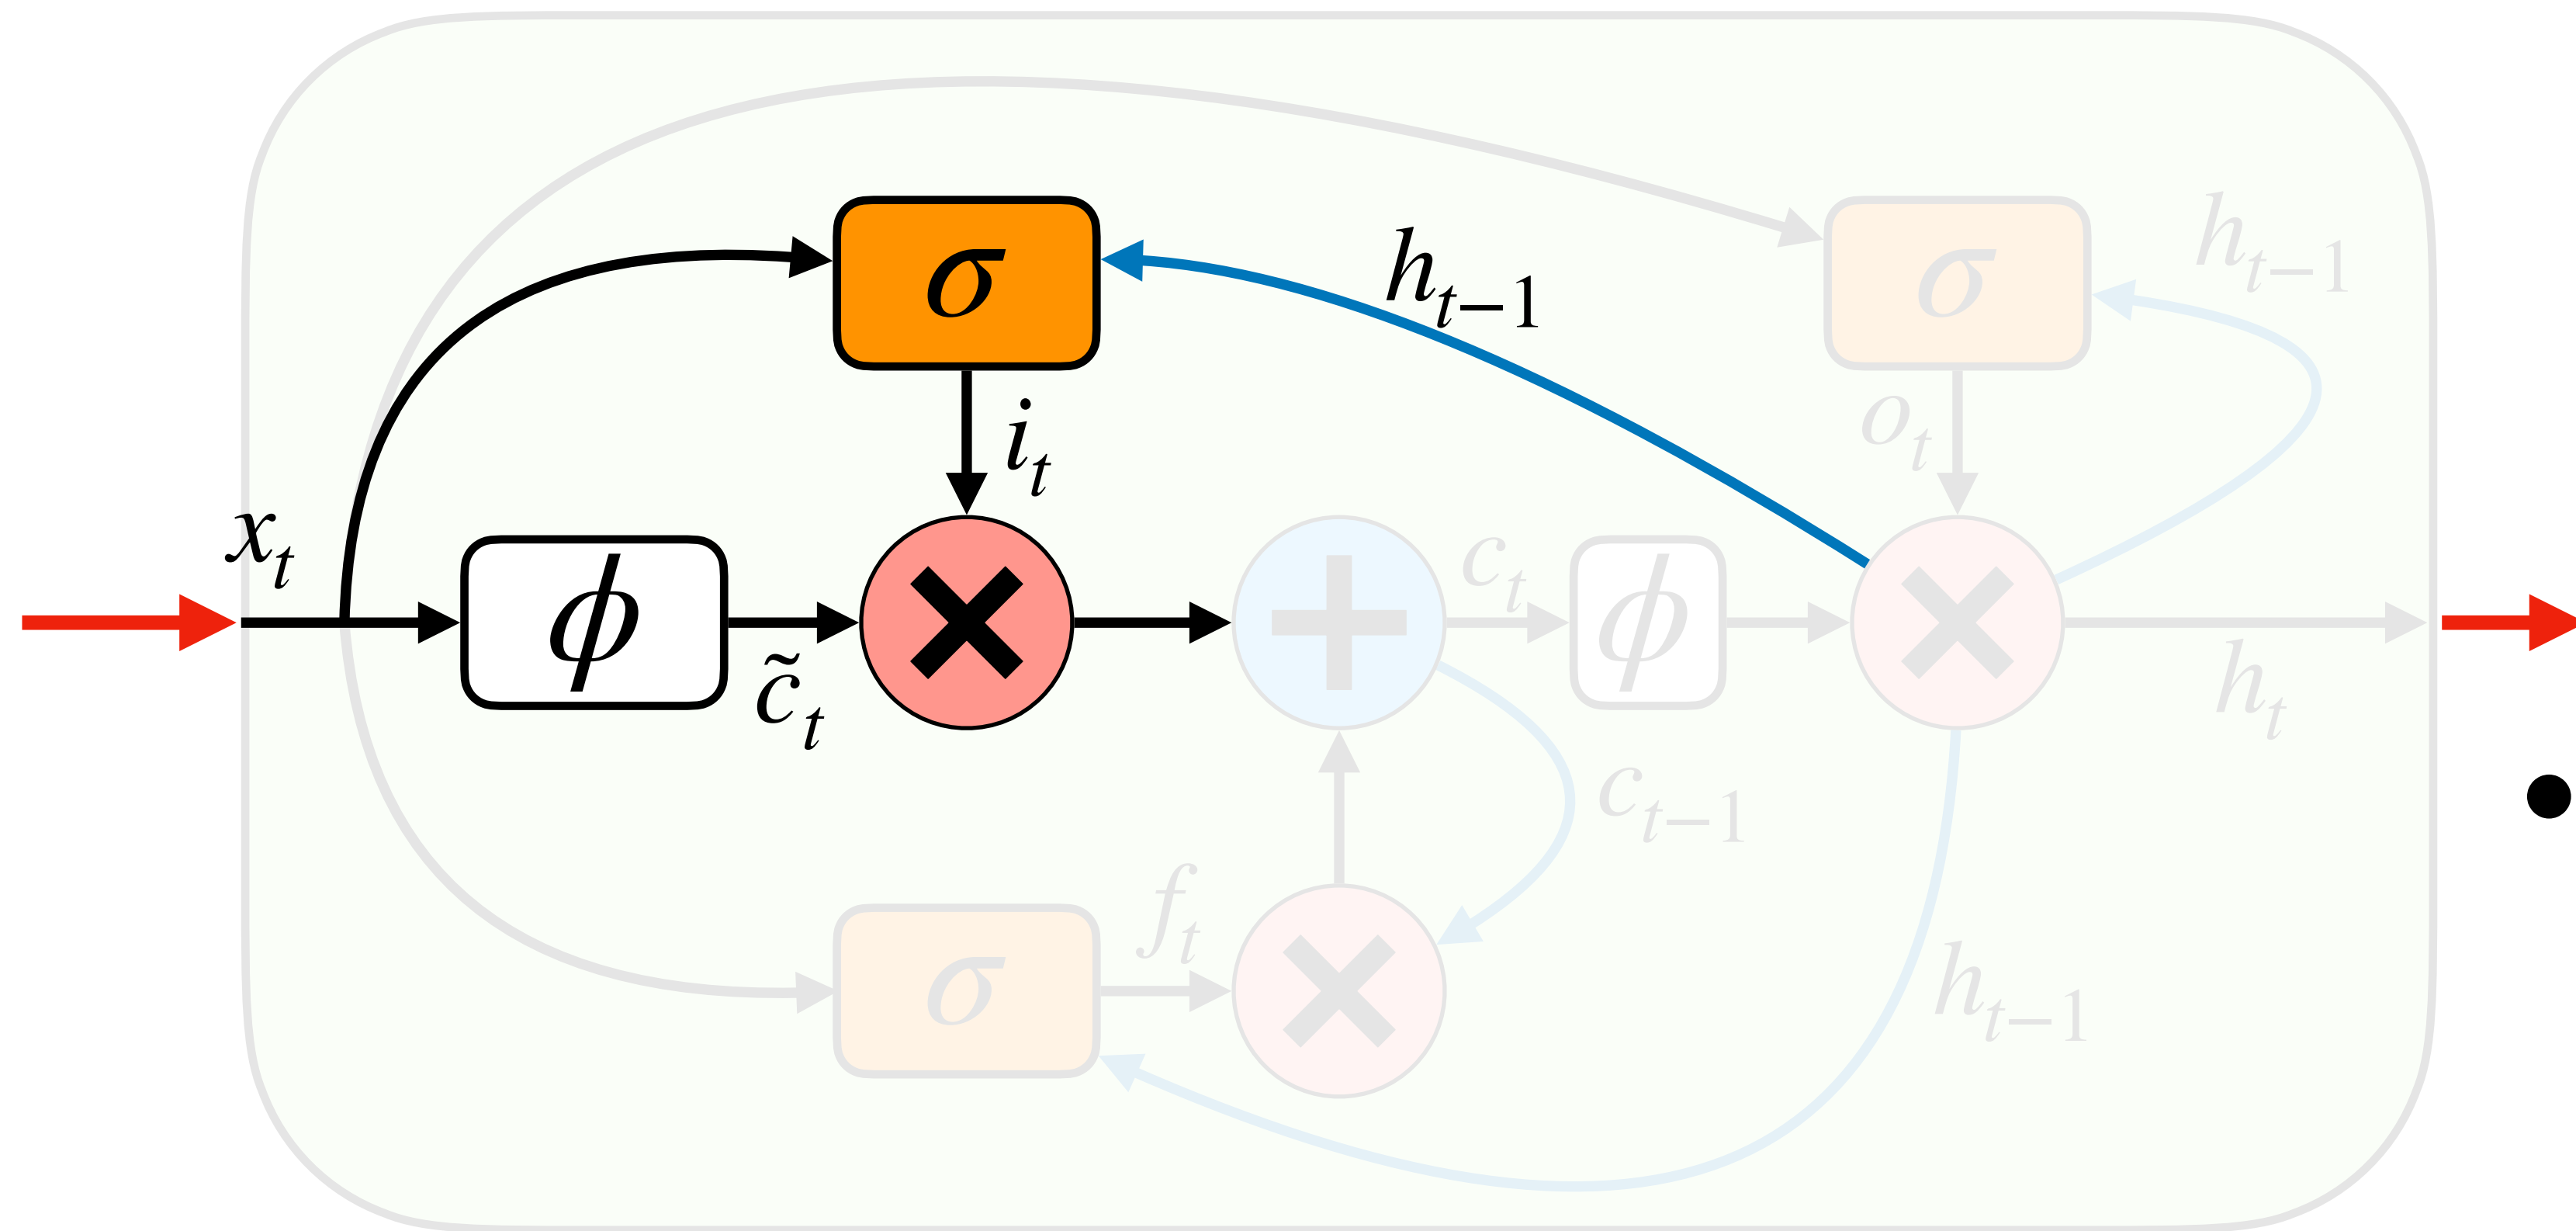 - 0 -> forget everything up to now

Input Gate

I went to **the lecture**

$$i_t = \sigma\big(W_{ix}x_t + W_{ih}h_{t-1} + b_i\big)$$

$$\tilde{c}_t = \phi\big(W_{cx}x_t + W_{ch}h_{t-1} + b_c\big)$$

$$c_t = i_t \times \tilde{c}_t + f_t \times c_{t-1}$$

- Input gate controls how new info is added to state memory

  - 0 -> ignore current time step

  - What might be ignored?

Output Gate

$$o_t = \sigma\big(W_{ox}x_t + W_{oh}h_{t-1} + b_o\big)$$

$$\tilde{c}_t = \phi\big(W_{cx}x_t + W_{ch}h_{t-1} + b_c\big)$$

$$c_t = i_t \times \tilde{c}_t + f_t \times c_{t-1}$$

$$h_t = o_t \times \phi(c_t)$$

- Output gate decides how the hidden state will influence gate values at next time step

Long Short Term Memory (LSTM)

Gates:

$$f_t = \sigma\big(W_{fx}x_t + W_{fh}h_{t-1} + b_f\big)$$

$$i_t = \sigma\big(W_{ix}x_t + W_{ih}h_{t-1} + b_i\big)$$
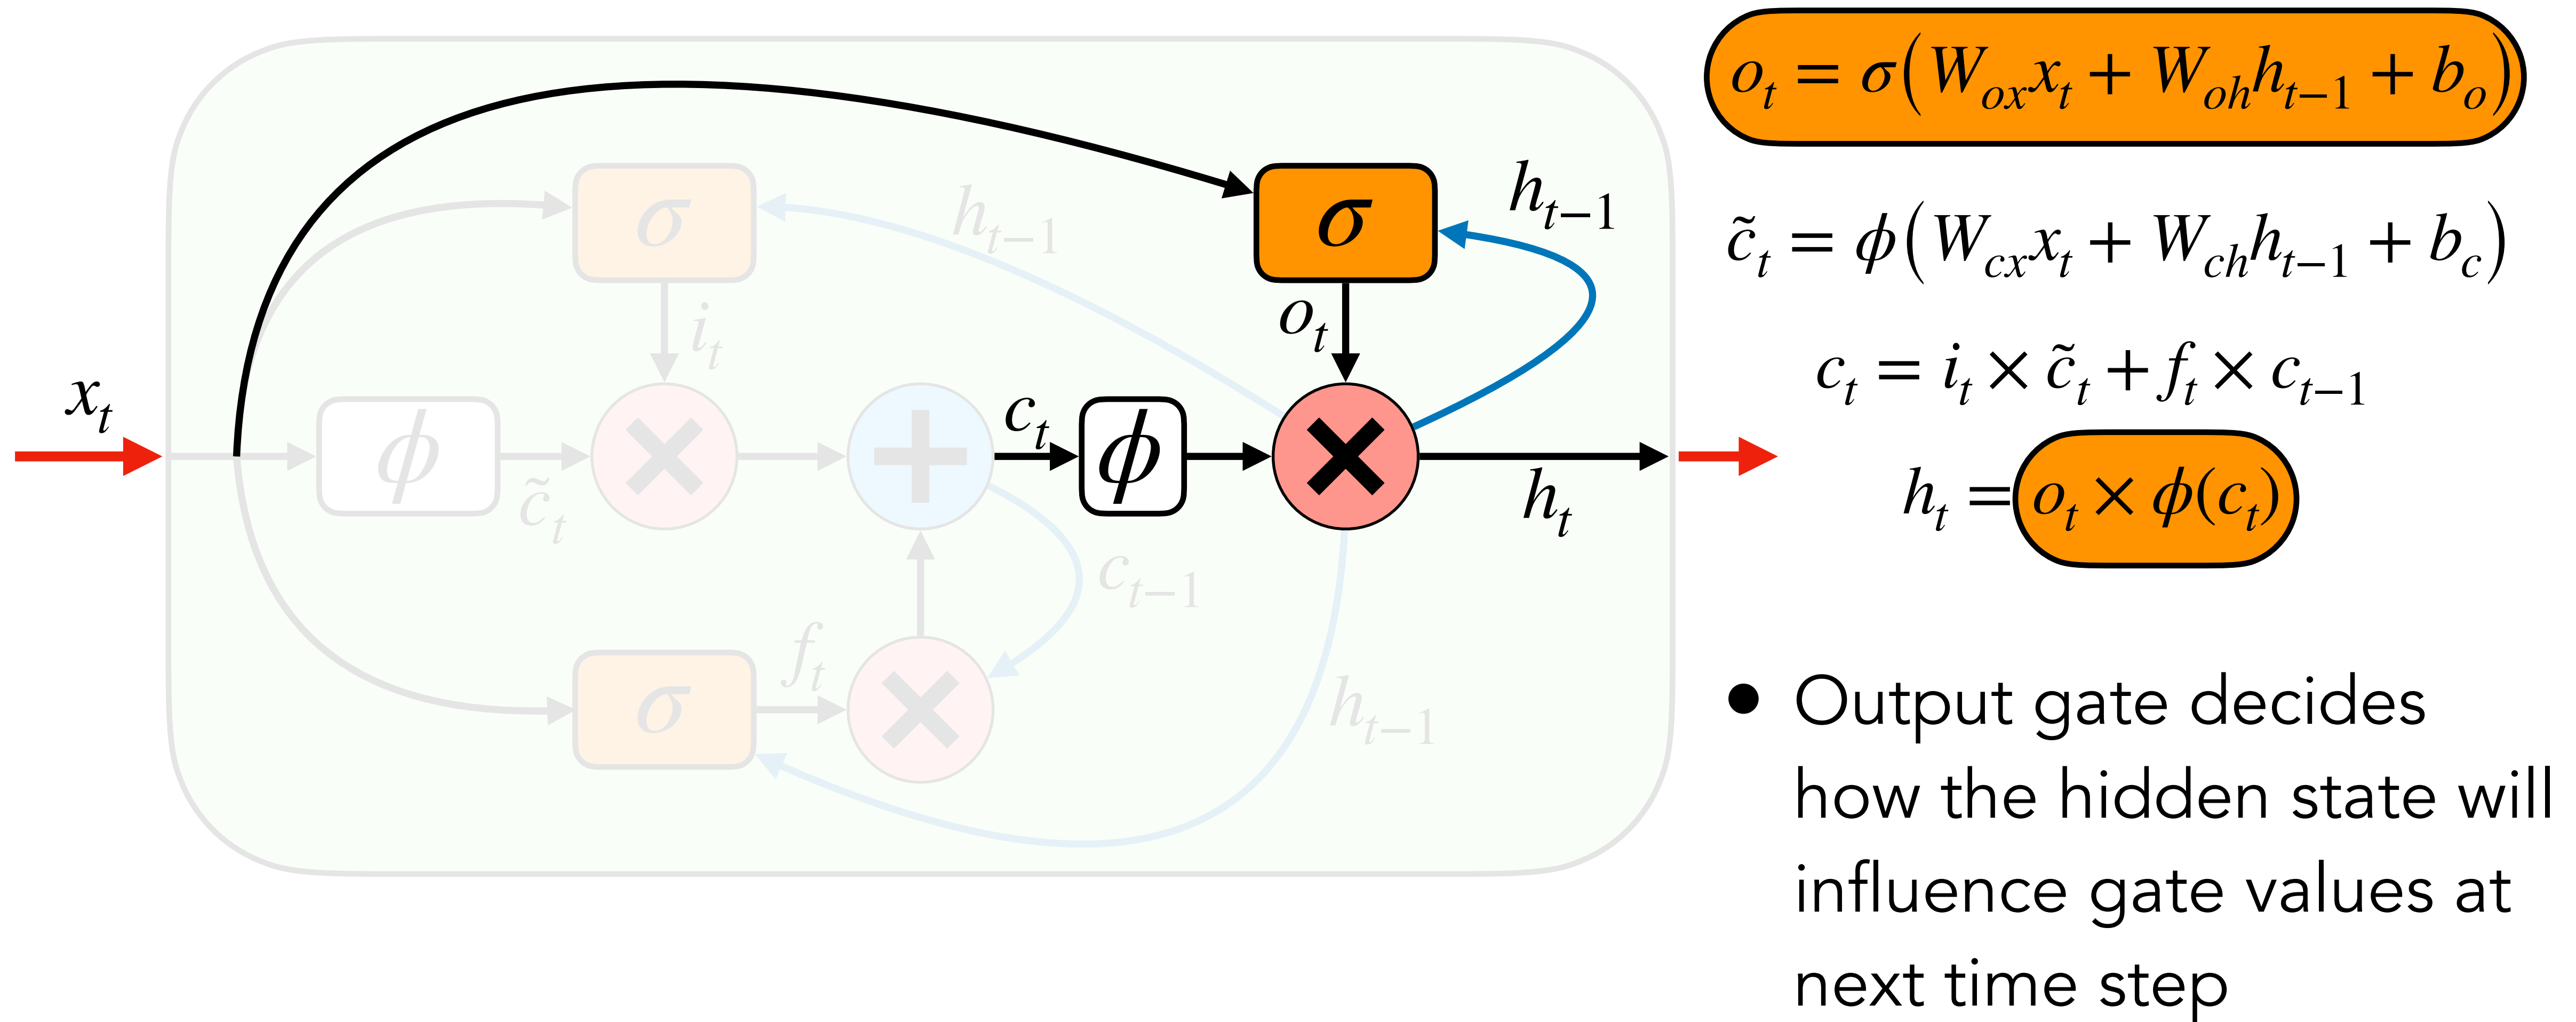
$$o_t = \sigma\big(W_{ox}x_t + W_{oh}h_{t-1} + b_o\big)$$

$$\tilde{c}_t = \phi\big(W_{cx}x_t + W_{ch}h_{t-1} + b_c\big)$$

$$c_t = i_t \times \tilde{c}_t + f_t \times c_{t-1}$$

$$h_t = o_t \times \phi(c_t)$$

Questions!

- For what type of input might the model learn to make the input gate be 0 ?

- What happens if the forget gate is 0?

- What happens if both the forget gate and input gate are 0 ?

- What happens if both the forget gate and input gate are 1 ?

Gates:

$$f_t = \sigma\left(W_{fx}x_t + W_{fh}h_{t-1} + b_f\right)$$

$$i_t = \sigma\left(W_{ix}x_t + W_{ih}h_{t-1} + b_i\right)$$

$$o_t = \sigma\left(W_{ox}x_t + W_{oh}h_{t-1} + b_o\right)$$

$$\tilde{c}_t = \phi\left(W_{cx}x_t + W_{ch}h_{t-1} + b_c\right)$$

$$c_t = i_t \times \tilde{c}_t + f_t \times c_{t-1}$$

$$h_t = o_t \times \phi(c_t)$$

Gated Recurrent Unit (GRU)

- Also uses gates to avoid dampening gradient signal every time step

$$h_t = (1 - \mathbf{z}) \odot h_{t-1} + \mathbf{z} \odot \mathbf{func}(x_t, h_{t-1}) \qquad h_t = h_{t-1} \odot \mathbf{f} + \mathbf{func}(x_t)$$
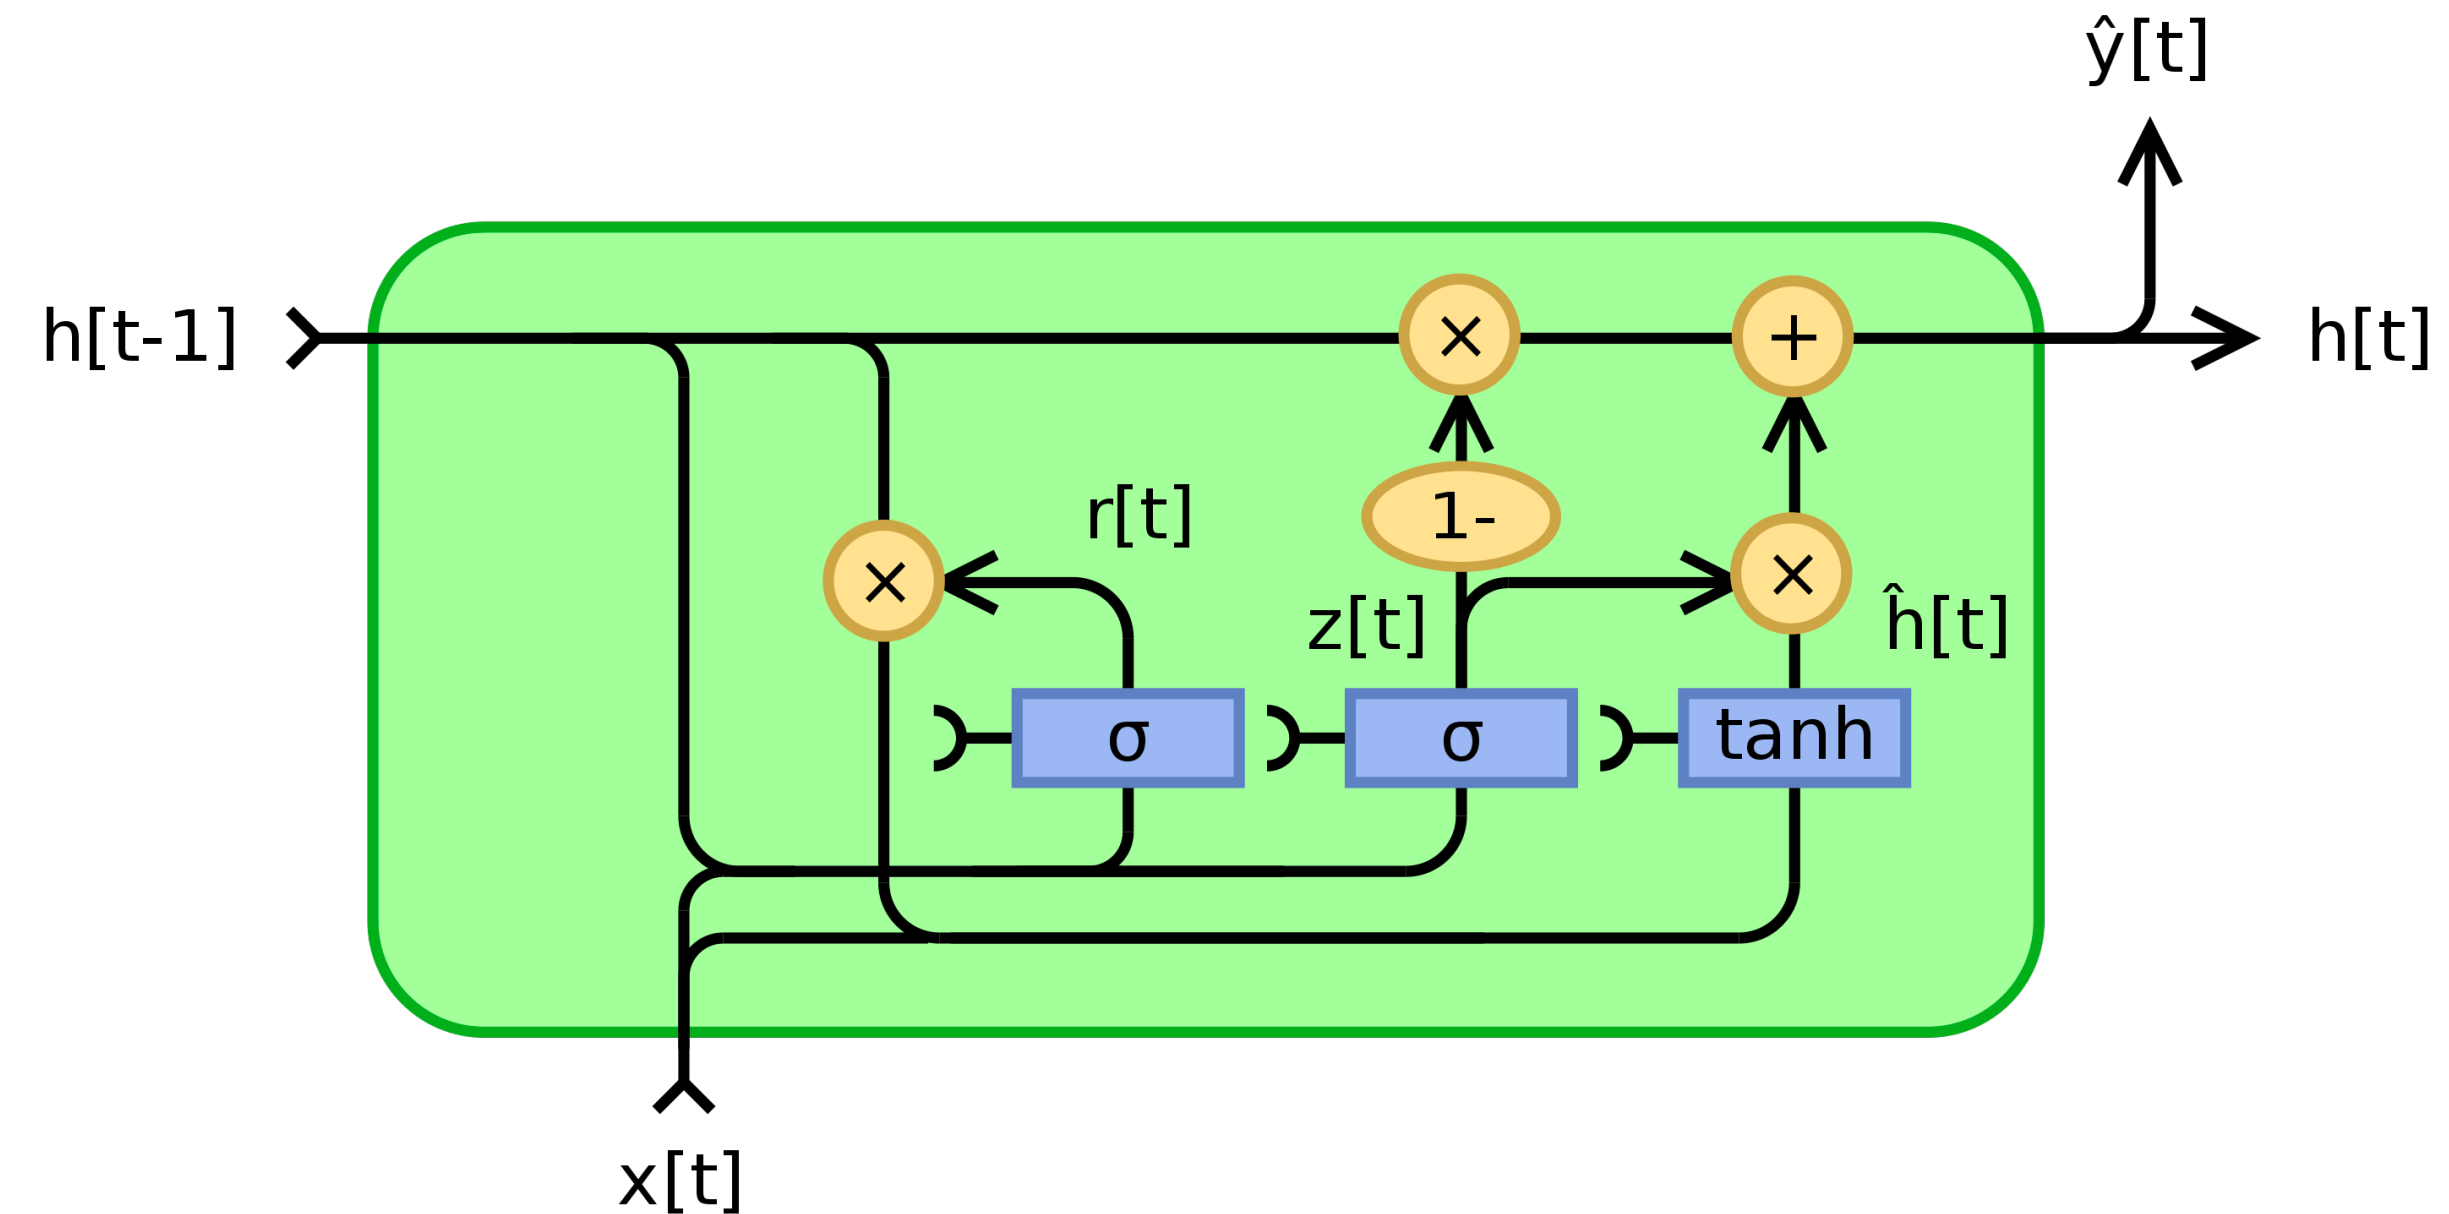
GRU   **LSTM**

- Works similarly to LSTM

  - Theoretically less powerful (find out why tomorrow!)

  - Typically faster to train and sometimes works better than LSTMs

Gated Recurrent Unit (GRU)

- **z** is update gate (used to update hidden state), **r** is reset gate (used to reset hidden state)

- The single hidden state and simpler update gate gives simpler mixing algorithm than in LSTMs

$$z_t = \sigma_g(W_z x_t + U_z h_{t-1} + b_z)$$
$$r_t = \sigma_g(W_r x_t + U_r h_{t-1} + b_r)$$
$$h_t = (1 - z_t) \circ h_{t-1} + z_t \circ \sigma_h(W_h x_t + U_h(r_t \circ h_{t-1}) + b_h)$$

Which is better?

Speech Signal Modeling

Music Modeling

Negative Loglikelihood

(Chung et al, 2014)

Question

What are the advantages of using LSTMs and GRUs?

Vanishing Gradients?

Recurrent Neural Networks	Long Short Term Memory

State maintained by hidden state feedback

$$h_t = \sigma\left(W_{hx}x_t + W_{hh}h_{t-1} + b_h\right)$$

Gradient systemically squashed by sigmoid

State maintained by cell value

$$c_t = i_t \times \tilde{c}_t + f_t \times c_{t-1}$$

Gradient set by value of forget gate

$$\frac{\partial c_t}{\partial c_{t-1}} = f_t$$

Can still vanish, but only if forget gate closes!

63

Question

What's a disadvantages of using a LSTM or GRU?

Question

What's a disadvantages of using a LSTM or GRU?

More parameters!

$$f_t = \sigma\big(W_{fx}x_t + W_{fh}h_{t-1} + b_f\big)$$

$$i_t = \sigma\big(W_{ix}x_t + W_{ih}h_{t-1} + b_i\big)$$

$$o_t = \sigma\big(W_{ox}x_t + W_{oh}h_{t-1} + b_o\big)$$

$$\tilde{c}_t = \phi\big(W_{cx}x_t + W_{ch}h_{t-1} + b_c\big)$$

$$c_t = i_t \times \tilde{c}_t + f_t \times c_{t-1}$$

$$h_t = o_t \times \phi(c_t)$$

$$z_t = \sigma\big(W_{zh}h_t + b_z\big)$$

$$h_t = \sigma\big(W_{hx}x_t + W_{hh}h_{t-1} + b_h\big)$$

Question

Could there be better architectures than GRUs and LSTMs?

Optimal Architectures?

MUT1:

$$z = \text{sigm}(W_{\text{xz}}x_t + b_{\text{z}})$$
$$r = \text{sigm}(W_{\text{xr}}x_t + W_{\text{hr}}h_t + b_{\text{r}})$$
$$h_{t+1} = \tanh(W_{\text{hh}}(r \odot h_t) + \tanh(x_t) + b_{\text{h}}) \odot z$$
$$+ \quad h_t \odot (1-z)$$

MUT2:

$$z = \text{sigm}(W_{\text{xz}}x_t + W_{\text{hz}}h_t + b_{\text{z}})$$
$$r = \text{sigm}(x_t + W_{\text{hr}}h_t + b_{\text{r}})$$
$$h_{t+1} = \tanh(W_{\text{hh}}(r \odot h_t) + W_{xh}x_t + b_{\text{h}}) \odot z$$
$$+ \quad h_t \odot (1-z)$$

MUT3:

$$z = \text{sigm}(W_{\text{xz}}x_t + W_{\text{hz}}\tanh(h_t) + b_{\text{z}})$$
$$r = \text{sigm}(W_{\text{xr}}x_t + W_{\text{hr}}h_t + b_{\text{r}})$$
$$h_{t+1} = \tanh(W_{\text{hh}}(r \odot h_t) + W_{xh}x_t + b_{\text{h}}) \odot z$$
$$+ \quad h_t \odot (1-z)$$

Arch.	Arith.	XML	PTB
Tanh	0.29493	0.32050	0.08782
LSTM	0.89228	0.42470	0.08912
LSTM-f	0.29292	0.23356	0.08808
LSTM-i	0.75109	0.41371	0.08662
LSTM-o	0.86747	0.42117	0.08933
LSTM-b	0.90163	0.44434	0.08952
GRU	0.89565	0.45963	0.09069
MUT1	**0.92135**	**0.47483**	0.08968
MUT2	0.89735	**0.47324**	0.09036
MUT3	0.90728	0.46478	**0.09161**

Arch.	5M-tst	10M-v	20M-v	20M-tst
Tanh	4.811	4.729	4.635	4.582 (97.7)
LSTM	4.699	4.511	4.437	4.399 (81.4)
LSTM-f	4.785	4.752	4.658	4.606 (100.8)
LSTM-i	4.755	4.558	4.480	4.444 (85.1)
LSTM-o	4.708	4.496	4.447	4.411 (82.3)
LSTM-b	4.698	4.437	4.423	**4.380 (79.83)**
GRU	4.684	4.554	4.559	4.519 (91.7)
MUT1	4.699	4.605	4.594	4.550 (94.6)
MUT2	4.707	4.539	4.538	4.503 (90.2)
MUT3	4.692	4.523	4.530	4.494 (89.47)

(Jozefowicz et al, 2015)

Recap

- Recurrent neural networks can **theoretically** learn to model an **unbounded context length**

  - no increase in model size because weights are shared across time steps

- Practically, however, **vanishing gradients** stop vanilla RNNs from learning useful **long-range dependencies**

- LSTMs and GRUs are variants of recurrent networks that mitigate the vanishing gradient problem

  - used for for **many sequence-to-sequence tasks (up next!)**

References

- Elman, J.L. (1990). Finding Structure in Time. *Cogn. Sci., 14*, 179-211.

- Schuster, M., & Paliwal, K. K. (1997). Bidirectional recurrent neural networks. *IEEE Transactions on Signal Processing, 45*(11), 2673–2681.

- Hochreiter, S., & Schmidhuber, J. (1997). Long Short-Term Memory. *Neural Computation, 9*, 1735-1780.

- Cho, K., Merrienboer, B.V., Gülçehre, Ç., Bahdanau, D., Bougares, F., Schwenk, H., & Bengio, Y. (2014). Learning Phrase Representations using RNN Encoder–Decoder for Statistical Machine Translation. *Conference on Empirical Methods in Natural Language Processing.*

- Greff, K., Srivastava, R.K., Koutník, J., Steunebrink, B.R., & Schmidhuber, J. (2015). LSTM: A Search Space Odyssey. *IEEE Transactions on Neural Networks and Learning Systems, 28*, 2222-2232.

References

- Bengio, Y., Ducharme, R., Vincent, P., & Janvin, C. (2003). A Neural Probabilistic Language Model. *Journal of machine learning research.*

- Elman, J.L. (1990). Finding Structure in Time. *Cogn. Sci., 14,* 179-211.